

Hugging Face in 4 Hours

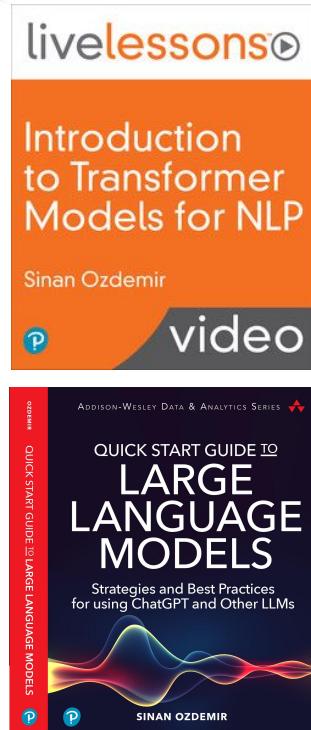


Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

Welcome!

My name is **Sinan Ozdemir** (in/sinan-ozdemir + [@prof_oz](https://twitter.com/prof_oz))



- Current **founder** of Loop Genius (using AI to help entrepreneurs get their first 100 customers)
- Current **lecturer** for O'Reilly and Pearson
- Founder of Kylie.ai (Funded by OpenAI Founder + Acquired)
- **Masters** in Theoretical Math from **Johns Hopkins**
- Former lecturer of Data Science at Johns Hopkins

Author of ML textbooks and online series, including

- [Quick Start Guide to LLMs](#)
- [Introduction to Transformer Models for NLP](#)



Hugging Face in 4 Hours



Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

Expectations for Today

We will spend most of our time together with my screen shared.

I will be showing off components of HuggingFace and code for using HuggingFace models, data, and APIs

Hugging Face in 4 Hours

Segment 1: Introduction to Hugging Face and Its Ecosystem



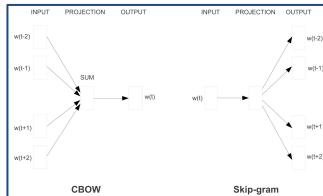
Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

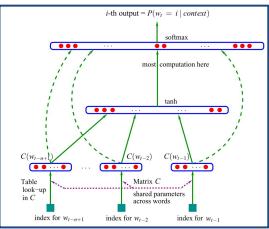
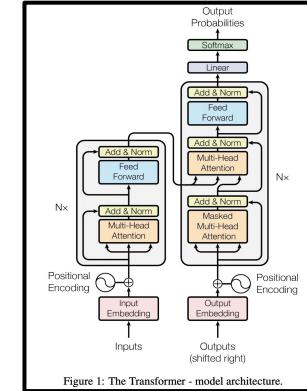
Introduction to transformer models and their significance in NLP

Brief History of Modern NLP

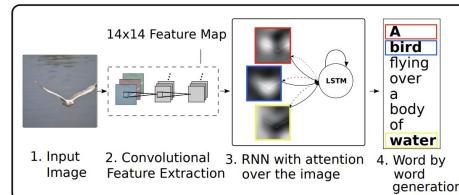
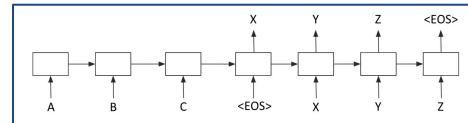
2001
Neural Language
Models



2014 - 2017
Seq2seq +
Attention



2013
encoding semantic
meaning with
Word2vec



2017 - Present
Transformers + Large
Language Models

2017 – Transformers

“Attention is all you need”

- Introduced the Transformer architecture
- A sequence to sequence model (takes text in and writes text back)
- The parent model of GPT3, BERT, T5, and many more

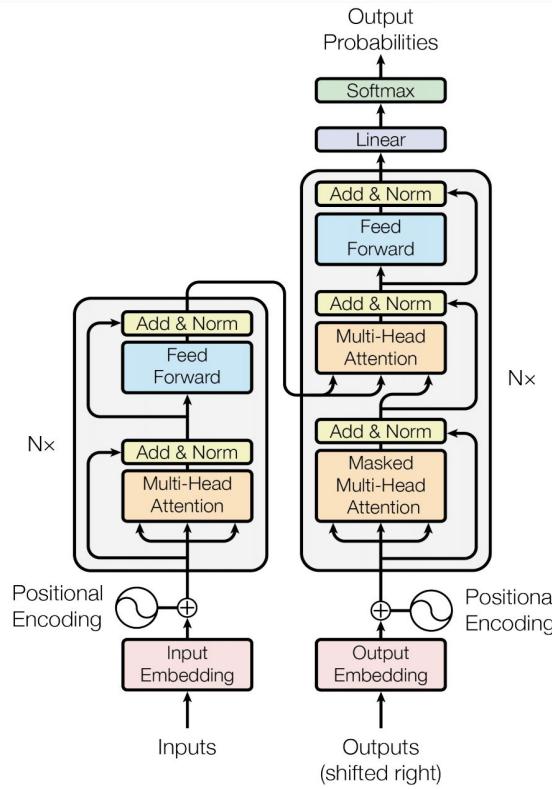


Figure 1: The Transformer - model architecture.

Auto-__ Language Models

Auto-regressive Models

Predict a future token (word) given either the past tokens or the future tokens but not both.

If you don't __ (forward prediction)

Auto-encoding Models

Learn representations of the entire sequence by predicting tokens given both the past and future tokens.

If you don't __ at the sign, you will get a ticket.

Using LLMs

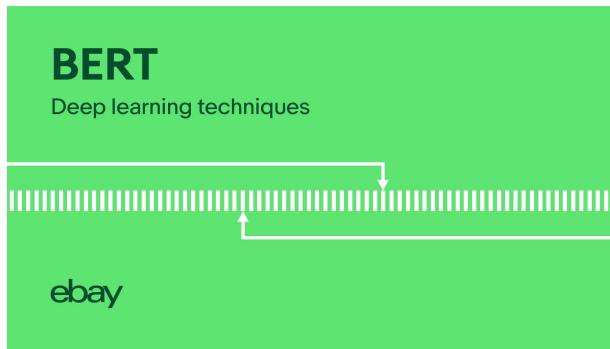
We can use LLMs in (generally) three ways:

1. **Encode** text into semantic vectors with little/no fine-tuning
 - a. Eg. Creating an information retrieval system using BERT vectors
2. Fine-tune a pre-trained LLM to perform a very specific task using **Transfer Learning**
 - a. Eg. Fine-tuning BERT to classify sequences with labels
3. Ask an LLM to solve a task it was pre-trained to solve or could intuit
 - a. Eg. **Prompting** GPT3 to write a blog post
 - b. Eg. **Prompting** T5 to perform language translation



Encoding Ebay's Recommendations with BERT

Ebay uses BERT to generate more relevant recommendations than traditional search techniques



The screenshot shows a search results page for "Michael Jordan" on eBay. At the top, there is a "Seed Item" image of a baseball card. Below it, two sections of recommendations are shown: "Recommendations without eBERT" and "Recommendations with eBERT". The "Recommendations with eBERT" section displays more relevant items, such as basketball cards from the 1993-94 Upper Deck Silver Top Prospects set and various Michael Jordan trading cards, including Fleer stickers and a "Feel The Game" basketball card. Each item listing includes a thumbnail image, the item title, price, and "Buy It Now" and "Add to cart" buttons.

Source:

<https://tech.ebayinc.com/engineering/how-ebay-created-a-language-model-with-three-billion-item-titles>

Overview of Hugging Face capabilities and community

Huggingface.co / models

Choose
the type
of model
you need



That's a lot

Hugging Face Search models, datasets, t

Models 487,674 Filter by name new Full-text search ↑↓ Sort: Trending

mistralai/Mixtral-8x7B-Instruct-v0.1
Text Generation • Updated Dec 15, 2023 • 1.21M • 2.49k

vikhyatk/moondream1
Updated 31 minutes ago • 192

InstantX/InstantID
Text-to-Image • Updated 8 days ago • 36.7k • 213

miqudev/miqu-1-70b
Updated 2 days ago • 134

stabilityai/stable-code-3b
Text Generation • Updated about 23 hours ago • 7.46k • 439

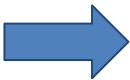
Tasks Libraries Datasets Languages Licenses Other Filter Tasks by name

Multimodal Feature Extraction Text-to-Image
Image-to-Text Image-to-Video
Text-to-Video Visual Question Answering
Document Question Answering
Graph Machine Learning Text-to-3D
Image-to-3D

Computer Vision Depth Estimation Image Classification
Object Detection Image Segmentation

Huggingface.co / model_page

Model tags including license



Model card for **allenai/OLMo-7B**

334 likes

Tags: Text Generation, Transformers, PyTorch, allenai/dolma, English, olmo, custom_code, Inference Endpoints

Arxiv IDs: arxiv:2402.00838, arxiv:2302.13971 | License: apache-2.0

Model card | Files | Community (9)

Train | Deploy | Use in Transformers

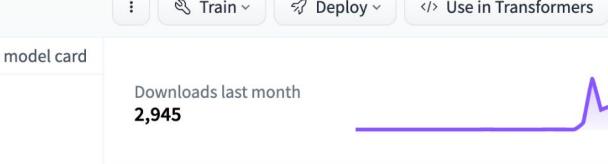


Model Card for OLMo 7B

General Information

OLMo is a series of **Open Language Models** designed to enable the science of language models. The OLMo models are trained on the

List ways to use the model



Text Generation

Model is too large to load onto the free Inference API. To try the model, launch it on [Inference Endpoints](#) instead.

Dataset used to train allenai/OLMo-7B

allenai/dolma
Updated 3 days ago • 468 likes



Collection including allenai/OLMo-7B

Huggingface.co / datasets

Choose
the type
of data
you need



 **Hugging Face** Models Datasets Spaces Posts Docs Pricing 

Tasks Sizes Sub-tasks Languages Licenses Other

Multimodal

 Feature Extraction  Text-to-Image

 Image-to-Text  Image-to-Video

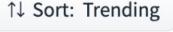
 Text-to-Video  Visual Question Answering

 Graph Machine Learning  Text-to-3D

 Image-to-3D

Computer Vision

 Depth Estimation  Image Classification

Datasets 102,276   Full-text search 


litagen/moe-speech
Updated about 1 hour ago • ↓ 69 • ❤ 158


fka/awesome-chatgpt-prompts
Viewer • Updated Mar 7, 2023 • ↓ 5.17k • ❤ 4.53k


PleIAs/French-PD-Newspapers
Viewer • Updated 4 days ago • ↓ 23 • ❤ 36


nampdn-ai/tiny-strange-textbooks
Viewer • Updated 3 days ago • ↓ 473 • ❤ 71

Huggingface.co / data_page

Data tags including license



Datasets: math-ai/AutoMathText like 53

Tasks: Text Generation, Question Answering

Languages: English

Size Categories: 10B < n < 100B

Tags: mathematical-reasoning, reasoning, finetuning, +2

License: cc-by-sa-4.0

Dataset card Files Community 1

Dataset Viewer (First 5GB)

Auto-converted to Parquet

API

Go to dataset viewer

Subset

Split

arxiv-0.50-to-1.00 (77.5k rows)

train (77.5k rows)

Search this dataset

url
string · lengths



<https://arxiv.org/abs/2105.10615>

title
string · lengths



Convergence directions of the...

abstract
string · lengths



The randomized Gauss-Seidel...

text
string · lengths



\section{Int...
problem is a
\section{Int...
methods for ...

Info on columns and showing rows



List ways to use data



Downloads last month

115

Use in Datasets library

Edit dataset card

⋮

Size of the auto-converted Parquet files (First 5GB per...
19.7 GB

Number of rows (First 5GB per split):
3,915,090



Show the dataset size

Huggingface.co / spaces

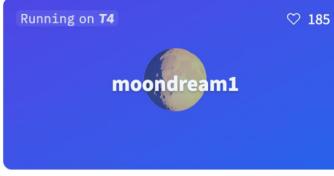
 **Spaces**

Discover amazing ML apps made by the community!

[Create new Space](#) or [learn more about Spaces.](#)

Search Spaces new Full-text search ↑↓ Sort: Trending

☆ Spaces of the week 🔥

 Depth Anything Web <small>Xenova 1 day ago</small>	 Live Vision <small>fffiloni 3 days ago</small>	 Running on A10G Internlm2 Math 7b <small>internlm 1 day ago</small>
 Running on ZERO InstantID <small>InstantX 4 days ago</small>	 Running on ZERO Compare Depth Models <small>merve 6 days ago</small>	 Running on T4 moondream1 <small>vikhaytk 6 days ago</small>

Huggingface.co / spaces

Spaces are a great way to see what's the latest and greatest in open-source

MoE-LLaVA: Mixture of Experts for Large Vision-Language Models

If you like our project, please give us a star ⭐ on Github for the latest update.

<https://github.com/PKU-YuanGroup/MoE-LLaVA>

The screenshot shows the MoE-LLaVA space interface. At the top, there's a logo featuring four anthropomorphic animal heads (lion, monkey, deer, fox) wearing graduation caps, with a computer monitor displaying a cat between them. Below the logo are GitHub, Code, Arxiv, and Stars buttons. A '1k' badge indicates the number of stars. A large dashed box labeled 'Input Image' contains the text 'Drop Image Here - or - Click to Upload'. Below this is an 'Examples' section with two rows. The first row shows a thumbnail of a person on a bicycle and the question 'What is unusual about this image?'. The second row shows a thumbnail of a wooden pier over water and the question 'What are the things I should be cautious about when I visit here?'. To the right, a card displays the question 'What is the title of this book?' above an image of the book 'QUICK START GUIDE TO LARGE LANGUAGE MODELS' by Sinan Ozdemir. The book cover features a colorful wavy graphic at the bottom. At the bottom of the card, the text reads 'Quick Start Guide to Large Language Models: Strategies and Best Practices for Using ChatGPT and Other LLMs'.

Huggingface.co / blog

Posts, articles, and discussions

[Everything](#) Community Guide Open Source Collab Partnerships
Research NLP Audio CV RL Ethics Diffusion Game Development
Time Series RLHF Case Studies



n the Hub

The Hallucinations Leaderboard, an Open Effort to Measure Hallucinations in Large Language Models

By pminervini · January 29, 2024 · guest

Community blog view posts all

Building autograd engine
tinytorch 03
By joey00072 · about 9 hou...

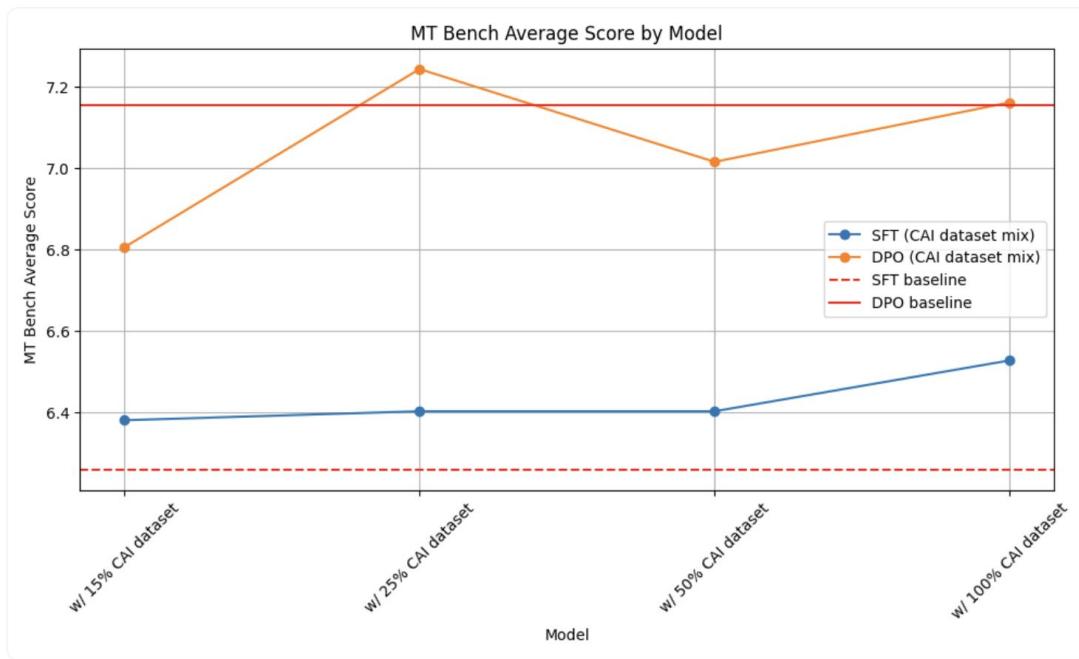
Building autograd engine
tinytorch 02
By joey00072 · about 9 hou...

Fine Tuning a LLM Using
Kubernetes with Intel® Xeon®
Scalable Processors
By dmsuehir · 7 days ago

Create a Web Interface for
your LLM in Python
By Alex1337 · 7 days ago

makeMoE: Implement a Sparse
Mixture of Experts Language
Model from Scratch

HF often has deep technical dives on things like Constitutional AI



We found training on the CAI dataset does not necessarily reduce helpfulness (i.e., paying the alignment tax). The SFT models obtained higher MT bench scores by training on

← Back to blog

Ethics and Society Newsletter #4: Bias in Text-to-Image Models

Both Giada
and Nathan
have been on
my show!



Published June 26, 2023

[Update on GitHub](#)



sasha
Sasha Luccioni



giadap
Giada Pistilli



nazneen
Nazneen Rajani



allendorf
Elizabeth Allendorf



irenesolaiman
Irene Solaiman



natolambert
Nathan Lambert



meg
Margaret Mitchell

TL;DR: We need better ways of evaluating bias in text-to-image models

Code Time!



Hugging Face in 4 Hours

Segment 2: Fine-tuning and Utilizing Pre-trained models

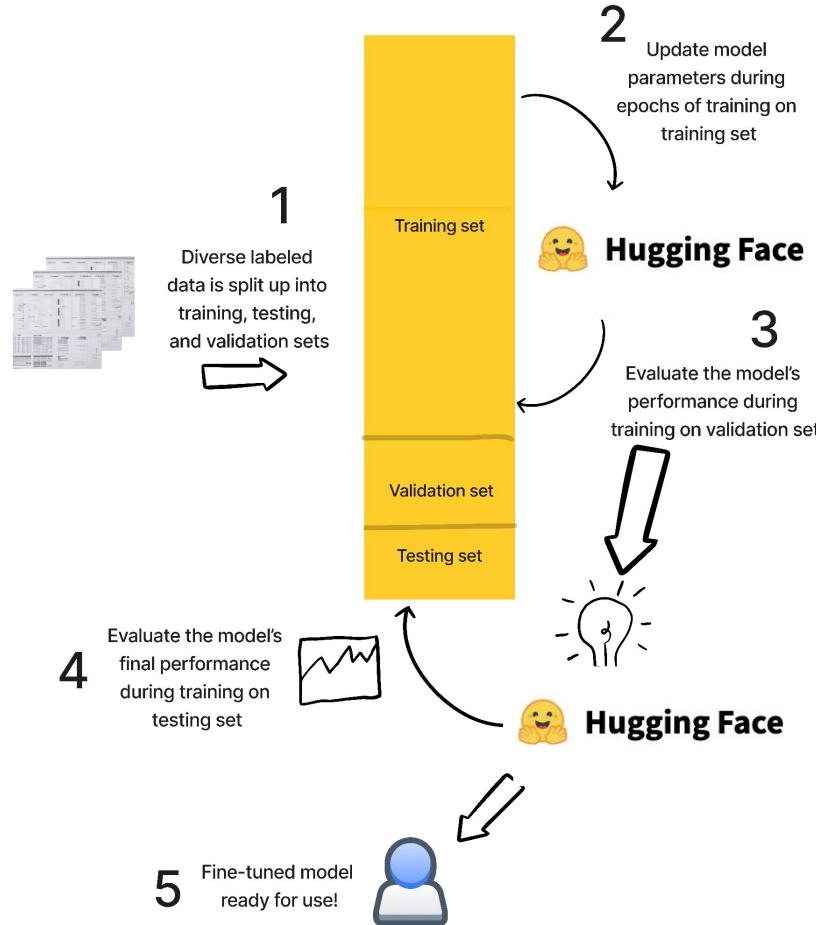


Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

Walkthrough of model fine-tuning process

Basic Fine-Tuning Process



Transfer Learning

Transfer Learning - A model trained for one task is reused as the starting point for a model for a second task.

1. Select a source model from a repository of models (like Huggingface)



2. Reuse and train the model for a second task using task-specific data

Transfer Learning with BERT



Tasks

- Fill-Mask
- Question Answering
- Summarization
- Table Question Answering
- Text Classification
- Text Generation
- Text2Text Generation
- Token Classification
- Translation
- Zero-Shot Classification
- Sentence Similarity
- + 16

Libraries

Models 35,367

distilgpt2

Text Generation • Updated May 21, 2021 • ↓ 33.2M • ❤ 39

bert-base-uncased

Fill-Mask • Updated May 18, 2021 • ↓ 16.2M • ❤ 125

facebook/bert-base-uncased

Text Classification • Updated Aug 5, 2021 • ↓ 11M • ❤ 5

Selecting a source model



Additional Task Layers

Pre-trained BERT

Training data for second task

Reusing and training model

BERT vs ChatGPT



Hugging Face



Search model



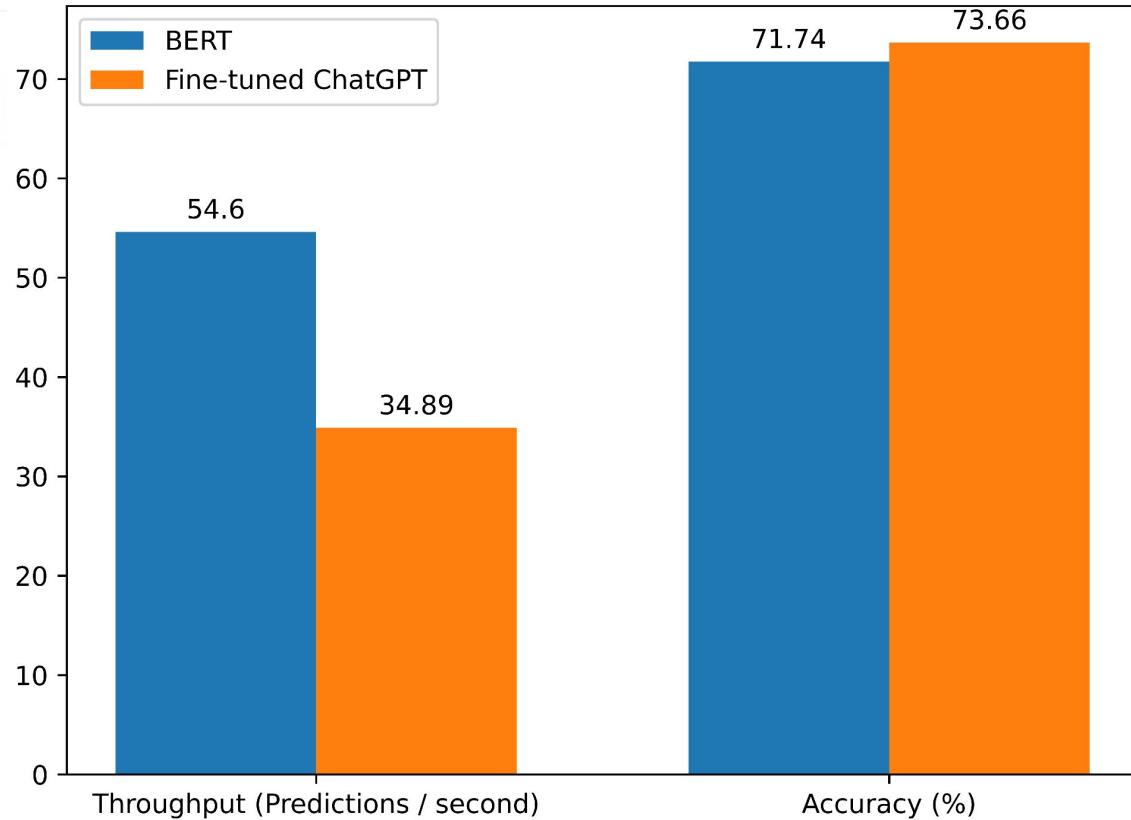
Datasets: app_reviews

Given a review, predict # stars

The BERT model has roughly 70M params and ChatGPT has ~175B

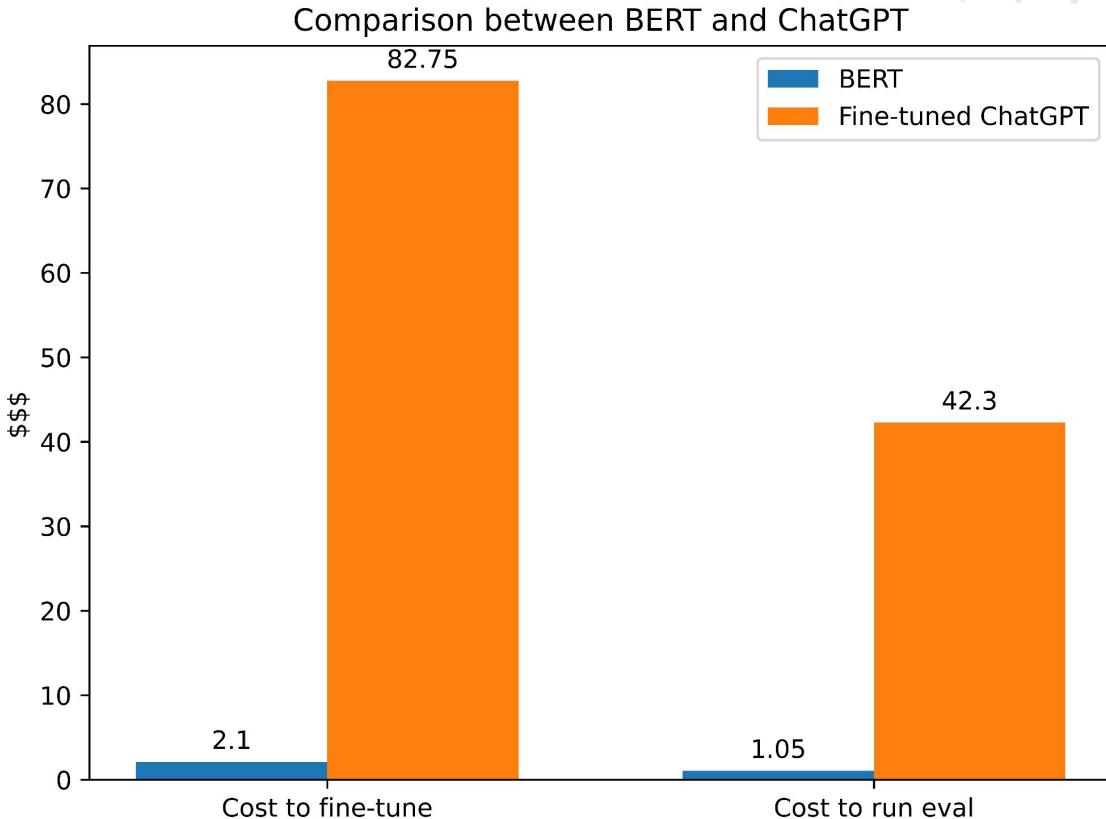
So BERT is ~2,500x smaller than ChatGPT but performances on par

Comparison between BERT and ChatGPT



BERT vs ChatGPT

BERT is also much cheaper / faster to train



Considering Open-source

Auto-encoding LLMs

Learns entire sequences by predicting tokens (words) given past and future context

If you don't __ at the sign, you will get a ticket.



cannot generate text but great for **classification**, **embedding + retrieval** tasks

Examples: **BERT**, XLNET, RoBERTa, sBERT

Auto-regressive LLMs

Predict a future token (word) given either past context or future context but not both.

If you don't __ mind? want? have?

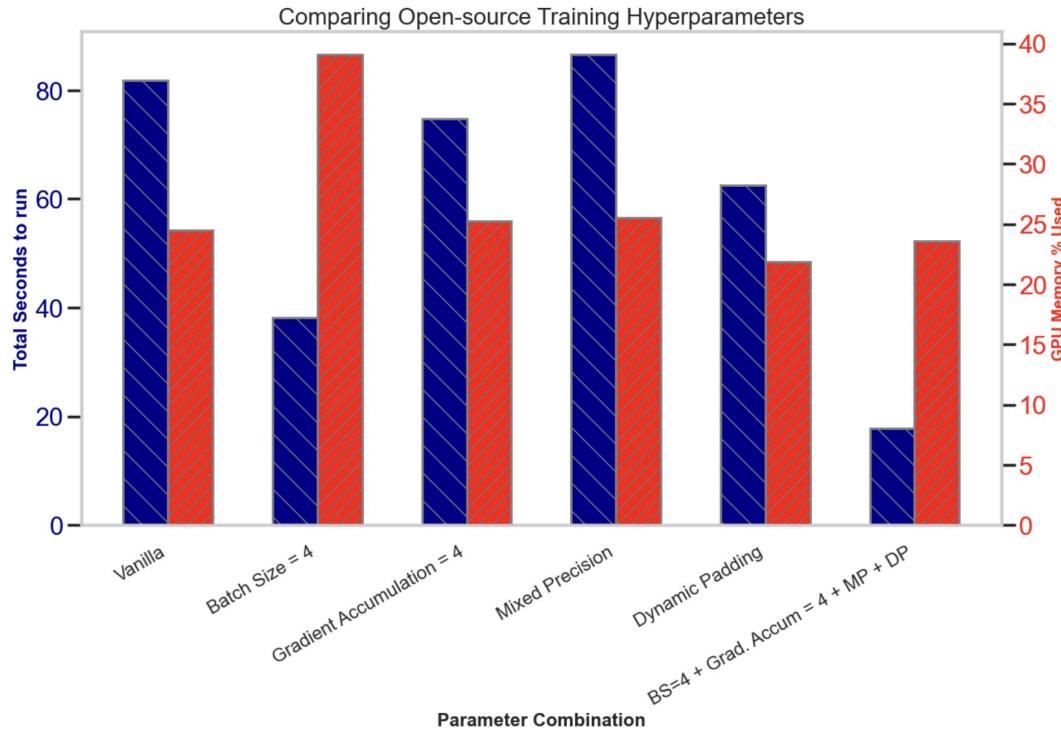
Capable of **generating text**, hence the term Generative LLMs but must be larger to read nearly as well as auto-encoding systems

Examples: **GPT** family, Llama family, Anthropic's Claude family, honestly most of the LLMs you see out there today

Optimizing Fine-tuning

Some smaller techniques (see more in my book or on my other lectures) can also be used to speed up training without consuming more memory.

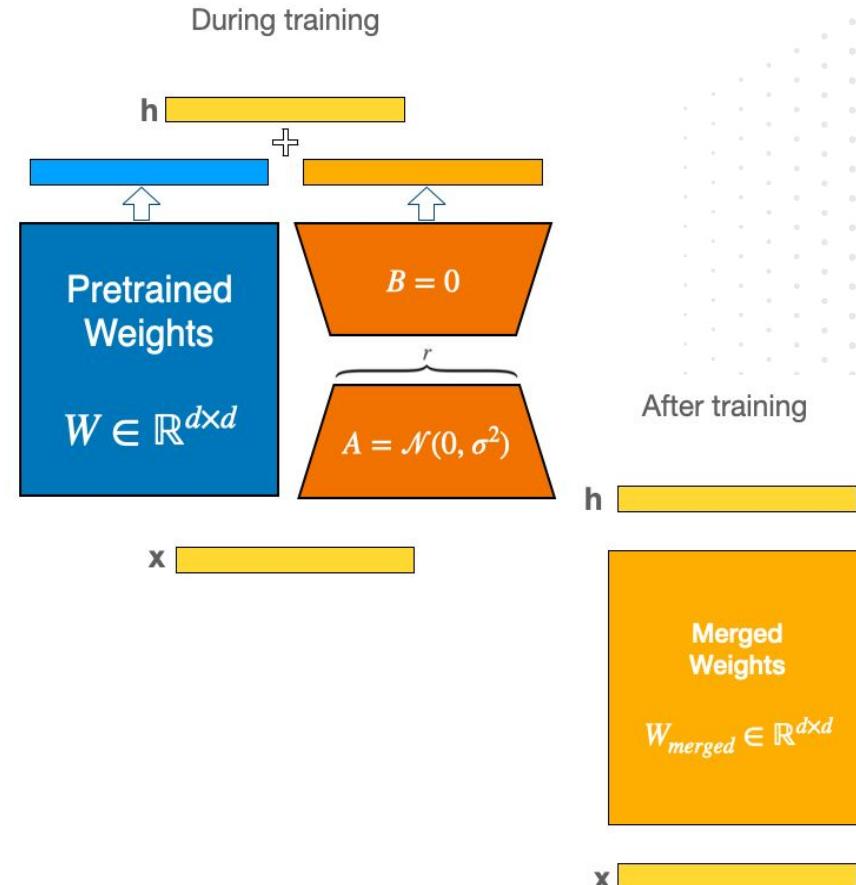
This graph shows a 4x speed up in training a classifier utilizing the same memory footprint



Advanced: PEFT to optimize memory

Parameter Efficient Fine-Tuning techniques like

LoRA (Low-Rank Adaptation) allow for training of larger models on smaller/single GPUs





Hugging Face

The AI community building the future.

27.3k followers

NYC + Paris

<https://huggingface.co/>

@huggingface

Verified

Pinned

transformers Public

🤗 Transformers: State-of-the-art Machine Learning for Pytorch, TensorFlow, and JAX.

Python 120k 23.9k

diffusers Public

🤗 Diffusers: State-of-the-art diffusion models for image and audio generation in PyTorch

Python 20.7k 4.3k

Exercise: Fine-tuning a sample model on a dataset

Code Time!



Code Time!



Hugging Face in 4 Hours

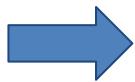
Segment 3: Deployment Strategies with Hugging Face



Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

ui.endpoints . huggingface.co



Deploy virtually
any model on
HF, even ones
you fine-tuned!

Deployed Endpoints + New Browse Catalog

Filter by name, provider, task... Status: All

aws-llama-2-7b-hf-3756 Scaled to Zero
text-generation • protected • aws • 1x Nvidia A10G
<https://jtjuckgwnk2jqym8.us-east-1.aws.endpoints.huggingface.cloud>

aws-llama-2-7b-chat-hf-6139 Scaled to Zero
text-generation • protected • aws • 1x Nvidia A10G
<https://sa0b44ky03zvbtos.us-east-1.aws.endpoints.huggingface.cloud>

distilbert-toxic-classifier Scaled to Zero
text-classification • protected • aws • Intel Ice Lake
<https://d2q5h5r3a1pkorfp.us-east-1.aws.endpoints.huggingface.cloud>

aws-t5-aligned-summaries-6363 Paused
text2text-generation • protected • aws • Intel Ice Lake
Endpoint has been paused

aws-flan-t5-small-1837 Paused
text2text-generation • protected • aws • Intel Ice Lake
Endpoint has been paused

aws-all-mppnet-base-v2-3472 Paused
sentence-embeddings • protected • aws • Intel Ice Lake
Endpoint has been paused

ui.endpoints . huggingface.co

distilbert-toxic-classifier

Running

|| Pause

Overview

Analytics

Usage & Cost

Logs

Settings

Endpoint URL

Need help?

https://d2q5h5r3a1pkorfp.us-east-1.aws.endpoints.huggingface.cloud

Model

Up-to-date

profoz/distilbert-toxic-classifier

Revision db64ff8

Configuration

Protected

Task

text-classification

Container Type

Default

Created Jan 25 by [profoz](#) • Last Edited Feb 5 at 9:16 AM by [proxy](#).

Instance

\$ 0.06/h while running



AWS



us-east-1

CPU · Intel Ice Lake · 1 vCPU · 2 GB

Scale-to-zero

after 15 minutes without activity



Task + endpoint
info



Compute info



Pearson

ui.endpoints . huggingface.co

Replicas

1 / 1

Total requests

180 req

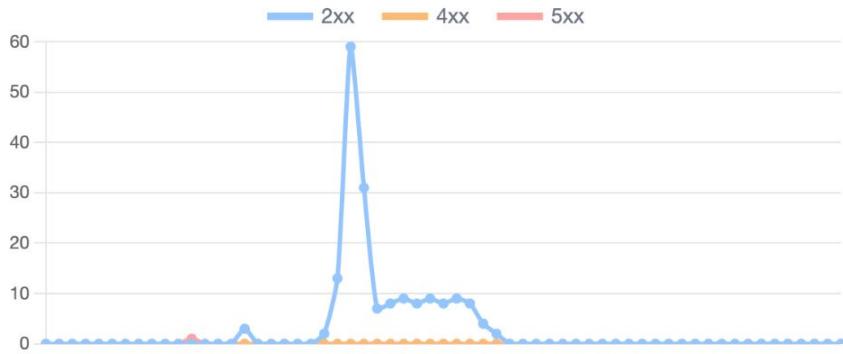
Median Latency

34.41 ms

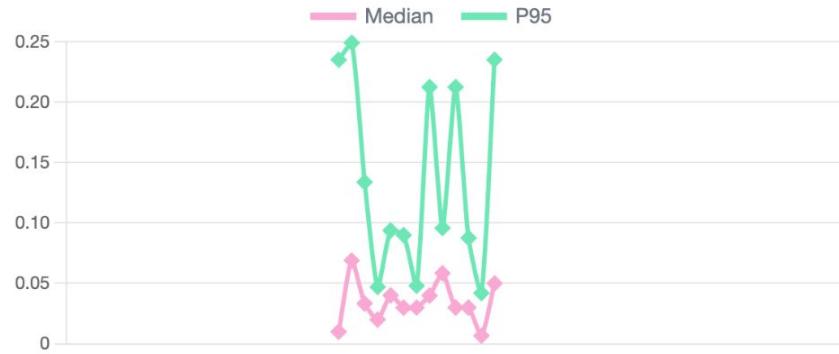
P95 Latency

137.08 ms

HTTP requests



Latency distribution (in seconds)



Basic Metrics

ui.endpoints . huggingface.co

API code to use the model

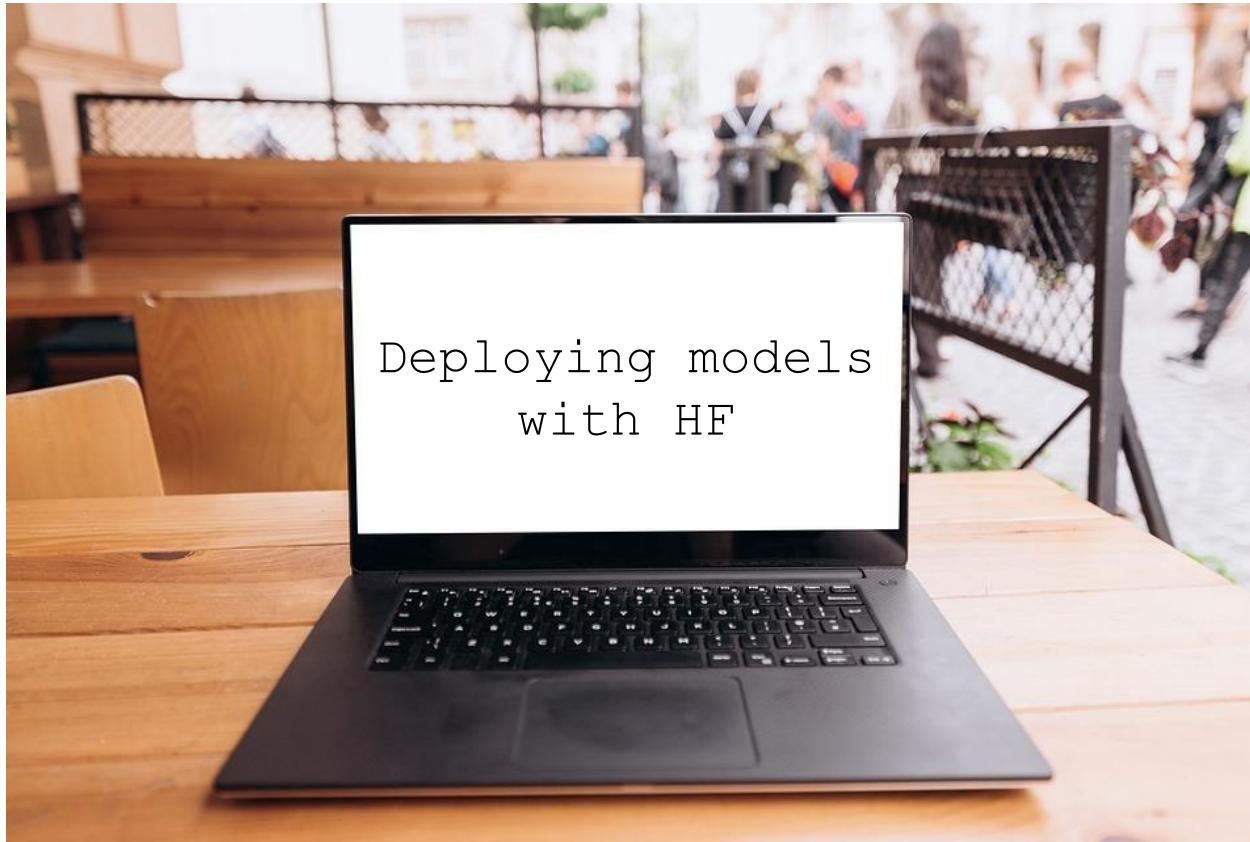
The screenshot shows the Hugging Face Model Playground interface. On the left, there's a sidebar with tabs for 'Playground', 'TEST' (which is selected), and 'API'. Below these are dropdowns for 'Token' (set to 'personal') and 'Parameters' (set to 'Top K' with 'number' input). Under 'Function to Apply', it says 'Default'. In the main area, there's a section titled 'Test your endpoint' with a lightning bolt icon. A text input field contains 'You're such a loser'. A 'Compute' button is below it. The results show a single item in a list: 'Toxic' with a score of '0.665'. At the bottom, there's a JSON output box containing the following code:

```
[{"label": "Toxic", "score": 0.6649302244186401}]
```

Simple
playground for
models

Exercise: Deploying a fine-tuned model on the Inference API

Code Time!



Chat UI



A chat interface using open source models, eg OpenAssistant or Llama. It is a SvelteKit app and it powers the [HuggingChat app on hf.co/chat](#).

0. [No Setup Deploy](#)
1. [Setup](#)
2. [Launch](#)
3. [Web Search](#)
4. [Text Embedding Models](#)
5. [Extra parameters](#)
6. [Deploying to a HF Space](#)
7. [Building](#)

Hugging Face in 4 Hours

Segment 4: Multimodal AI and Community Insights



Sinan Ozdemir

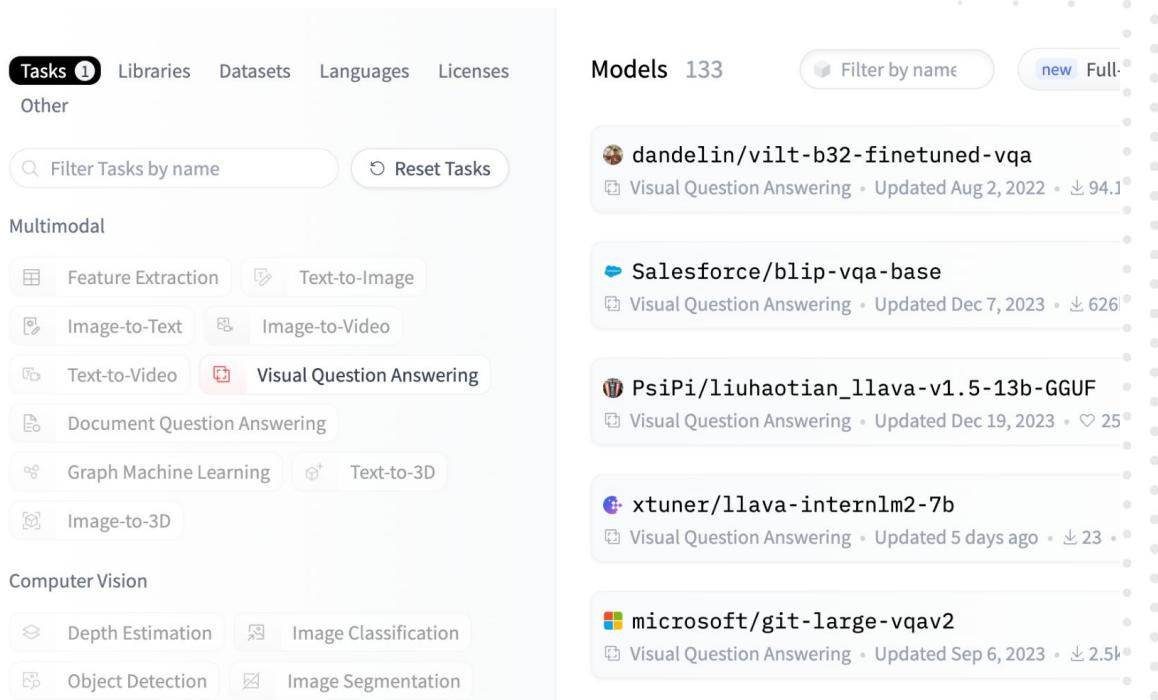
Data Scientist, Entrepreneur,
Author, Lecturer

Engaging with multimodality in AI: text, image, and audio processing

Multimodal Models

Multimodal models employ a
mix of data modalities

Computer Vision models are
strictly working with images



The screenshot shows the Hugging Face Model Hub interface. At the top, there are tabs for 'Tasks' (selected), 'Libraries', 'Datasets', 'Languages', and 'Licenses'. Below the tabs is a search bar labeled 'Filter Tasks by name' and a 'Reset Tasks' button. The main content is divided into two sections: 'Multimodal' and 'Computer Vision'. The 'Multimodal' section contains cards for Feature Extraction, Text-to-Image, Image-to-Text, Image-to-Video, Text-to-Video, Visual Question Answering, Document Question Answering, Graph Machine Learning, Text-to-3D, and Image-to-3D. The 'Computer Vision' section contains cards for Depth Estimation, Image Classification, Object Detection, and Image Segmentation. To the right, a list of 133 models is shown, each with a profile picture, name, description, and update date. The first few models listed are: danelin/vilt-b32-finetuned-vqa, Salesforce/blip-vqa-base, PsiPi/liuhaojian_llava-v1.5-13b-GGUF, xtuner/llava-internlm2-7b, and microsoft/git-large-vqav2.

Model	Description	Last Updated	Score
danelin/vilt-b32-finetuned-vqa	Visual Question Answering	Aug 2, 2022	94.1
Salesforce/blip-vqa-base	Visual Question Answering	Dec 7, 2023	626
PsiPi/liuhaojian_llava-v1.5-13b-GGUF	Visual Question Answering	Dec 19, 2023	25
xtuner/llava-internlm2-7b	Visual Question Answering	5 days ago	23
microsoft/git-large-vqav2	Visual Question Answering	Sep 6, 2023	2.51

Multimodal Models

A common multimodal task is **Visual Question/Answer**

Given an image and a question, answer the question (usually only with a word or two)

⚡ Inference API ⓘ

Visual Question Answering

Example 2 ▾



are we outside or inside?

Compute

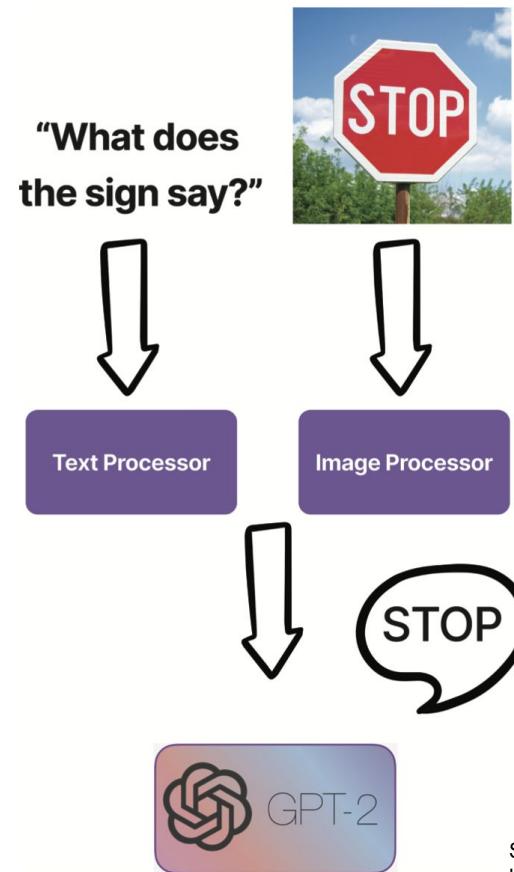
Computation time on cpu: 0.207 s

outside

0.990

Multimodal Models

You can build multimodal architectures using open-source components from HuggingFace



Source: Quick Start Guide to LLMs
by Sinan Ozdemir

Code Time!



Leveraging the community for project collaboration and advancement

Considering Open-source

Collaboration

Within org:

Setting up channels of communication between Data teams <> Product teams <> Marketing e.g. for **faster iteration** on meaningful features

Outside of org:

Sharing open-source models/tools is a chance to build **community** - a low-cost marketing strategy

Privacy / Security

Nothing new here, no need to send data to a 3rd party provider like **OpenAI** who have already shown a record of **data leaks** in their (relatively) short time in the limelight.

Ownership

Ownership of models and data provides an opportunity for organizations to get **more hands-on** with their ML use-cases by labeling data and **collecting feedback** from users.



Huggingface.co / blog

Posts, articles, and discussions

[Everything](#) Community Guide Open Source Collab Partnerships
Research NLP Audio CV RL Ethics Diffusion Game Development
Time Series RLHF Case Studies



n the Hub

The Hallucinations Leaderboard, an Open Effort to Measure Hallucinations in Large Language Models

By pminervini · January 29, 2024 · guest

Community blog view posts all

Building autograd engine
tinytorch 03
By joey00072 · about 9 hou...

Building autograd engine
tinytorch 02
By joey00072 · about 9 hou...

Fine Tuning a LLM Using
Kubernetes with Intel® Xeon®
Scalable Processors
By dmsuehir · 7 days ago

Create a Web Interface for
your LLM in Python
By Alex1337 · 7 days ago

makeMoE: Implement a Sparse
Mixture of Experts Language
Model from Scratch

Sinan Ozdemir's Framework for prototyping with LLMs with a mind for production

Sinan's LLM Framework

1. Define Inputs and Outputs

- Identify and document the specific inputs and outputs for your LLM application.
- Example: Given a user's taste and a list of book descriptions, the model should output a ranked list of book recommendations with reasons.
- Remember, requirements might change during testing or in different contexts.

2. Define Success/Failure States

- Clearly define what constitutes a success or a failure for your model.
- Example of success: The model should return at least 3 recommendations that match the given book list with a rationale for each.
- Example of failure: The model doesn't provide 3 recommendations, or the suggestions aren't from the given list.
- Failures are binary and don't reflect the quality of output, instead indicating whether the model meets the basic requirements.



Sinan's LLM Framework

3. Consider Potential Bias

- Examine if the model's outputs can be influenced by subjective bias or unnecessary information.
- Example: The model might utilize past knowledge or context about the books, leading to bias. Ensure it's "staying on script" and relying on the input given.

4. Create Comprehensive Examples (to be used as few-shot later)

- Develop at least two detailed examples for training (few-shot) or testing.
- Example: real list of wines from a dataset, etc
- This step helps to classify the model's knowledge requirement (Class A, B, or C).

Sinan's LLM Framework

5. Determine the Model's Knowledge Requirement

- Assess if the model has the necessary information to perform the task.
 - Class A: The model has all the required information encoded.
 - Class B: The model mostly has the necessary information but lacks specific details or updated data.
 - Class C: The model lacks the majority of required knowledge and needs extensive training.

6. Write an MVP (Minimum Viable Product) Prompt

- Create various versions of a prompt and experiment with them in the model's playground. This helps to refine the prompts and assess the model's knowledge requirement.

7. Iterate on Prompt Techniques and Parameters

- Adjust the parameters like temperature and top-p to refine the model's responses.



Sinan's LLM Framework

8. Evaluate and Plan for Scale/Production/Cost/Testing

- Assess the performance of the model, including its computational demands, and plan for potential scaling and production deployment.
- Also, consider the cost of deployment, which includes financial costs (like cloud resources and potential fine-tuning) and resource costs (like time and personnel for testing and maintenance).

9. Prototyping and Iteration

- Create a basic version of the model using tools like Streamlit for quick testing and user feedback.
- Iterate on the model by refining the prompts, adjusting parameters, and fine-tuning the model based on feedback.



Sinan's LLM Framework

10. Labeling Data and Fine-tuning

- Plan for potential data labeling and fine-tuning. This includes considering the cost and time required for these steps.
- Remember, fine-tuning not only requires labeled data but also extensive computational resources, which can add to the overall cost.

11. Evaluation

- Consistently evaluate the model's performance using relevant metrics like semantic similarity, precision, recall, etc. These evaluations will guide the iterations and improvements.

The above framework is not exhaustive but provides a good starting point for designing applications with LLMs like ChatGPT. Each application will have unique needs and constraints, so this framework should be adapted accordingly.



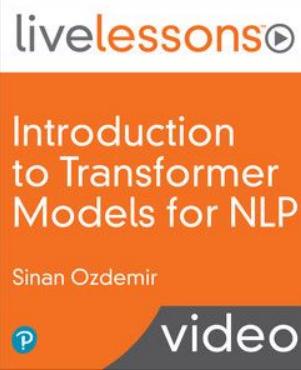
Summary + Next Steps

- The invention of the Transformer in 2017 revitalized of the field of NLP and an explosion of Large Language Models
- There are many types of LLMs with pros/cons and knowing which to use and how to use it makes all the difference
- LLMs are not perfect and **will** eventually produce untrue and harmful statements if left unchecked
- Reinforcement Learning can further align LLMs
- Attention seems to be (mostly) all we need.. for now

Summary + Next Steps

- Libraries like Streamlit help fasttrack prototypes and give you the ability to share them for free on Hugging Face
- Knowing which metrics are best for evaluation can make all the difference
- Building prototypes off of a framework and using future-proof techniques like few-shot prompting and chain-of-thought reasoning help us build faster and with more confidence

Summary + Next Steps

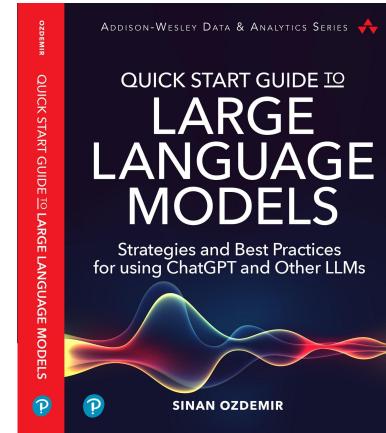


A comprehensive introduction to LLMs + Transformers

<https://learning.oreilly.com/videos/introduction-to-transformer/9780137923717>

Check out my live trainings for more in depth content!

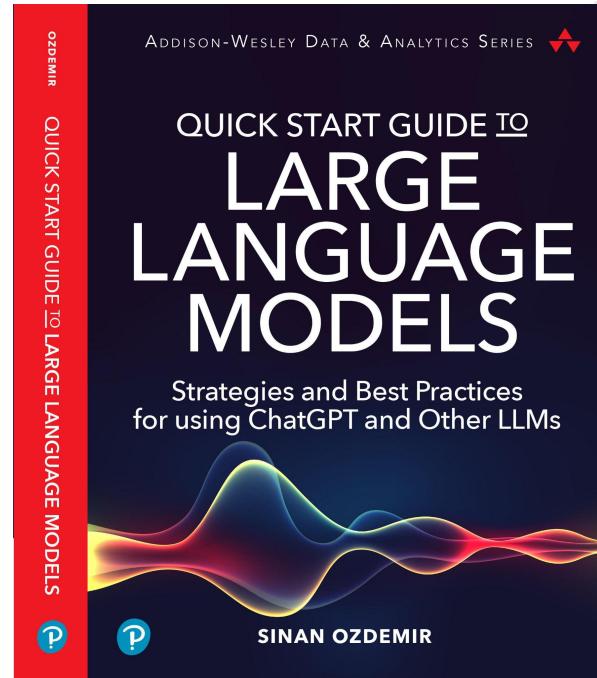
<https://learning.oreilly.com/search/?q=Sinan%20Ozdemir&type=live-event-series>



Thank you! / Final Q/A

Most of these examples were based off of my new book on LLMs, usually top 10 in many categories on Amazon including NLP

<https://a.co/d/fZsOWxd>



Hugging Face in 4 Hours

Thank you!



Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer