Large Language Models and ChatGPT in 3 Weeks

Week 3 - Building Viable Prototypes with LLMs and GPT

Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer

Pearson

Welcome!

My name is **Sinan Ozdemir** (in/sinan-ozdemir + @prof_oz)

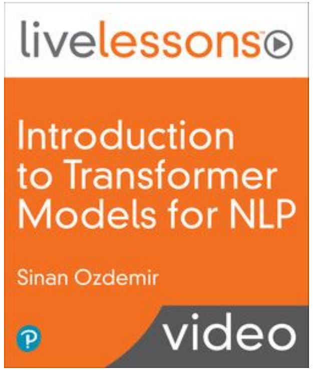

livelessons

Introduction to Transformer Models for NLP

Sinan Ozdemir

video

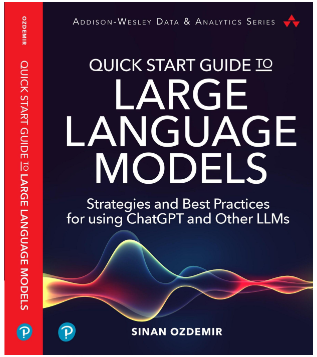

Addison-Wesley Data & Analytics Series

QUICK START GUIDE TO LARGE LANGUAGE MODELS

Strategies and Best Practices for using ChatGPT and Other LLMs

SINAN OZDEMIR

OZDEMIR QUICK START GUIDE TO LARGE LANGUAGE MODELS

Pearson

- Current **founder** of Loop Genius (using AI to help entrepreneurs get their first 100 customers)
- Current **lecturer** for O'Reilly and Pearson
- Founder of Kylie.ai (Funded by OpenAI Founder + Acquired)
- **Masters** in Theoretical Math from **Johns Hopkins**
- Former lecturer of Data Science at Johns Hopkins

Author of ML textbooks and online series, including

- [The Principles of Data Science]()
- [Introduction to Transformer Models for NLP]()
- [Quick Start Guide to LLMs]() (Top 10 in NLP on Amazon)

Quick Recap

Brief History of Modern NLP

2001
Neural Language Models

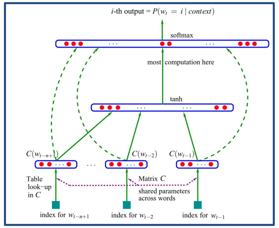

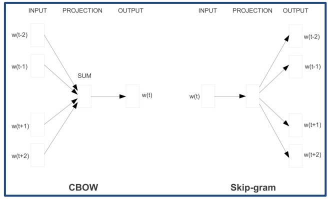

2013
encoding semantic meaning with Word2vec

2014 - 2017
Seq2seq + Attention

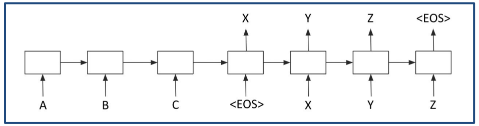

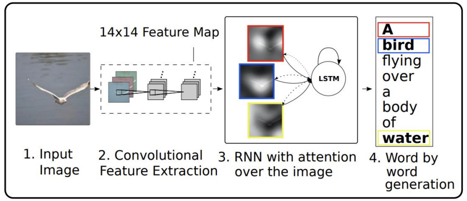
14x14 Feature Map
LSTM
A
bird
flying
over
a
body
of
water
1. Input Image
2. Convolutional Feature Extraction
3. RNN with attention over the image
4. Word by word generation

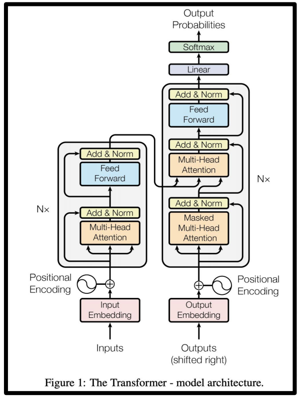

Figure 1: The Transformer - model architecture.

2017 - Present
Transformers + Large Language Models

Bengio et al. https://www.jmlr.org/papers/volume3/bengio03a/bengio03a.pdf
Mikolov et al. https://arxiv.org/abs/1301.3781
Xu et al. http://proceedings.mlr.press/v37/xuc15.pdf
https://papers.nips.cc/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf

Auto-__ Language Model Use Cases

Auto-regressive Models

1. Predicting next word in a sentence (auto-complete)
2. Natural Language Generation (NLG)
3. GPT Family

Auto-encoding Models

1. Comprehensive understanding and encoding of entire sequences of tokens
2. Natural Language Understanding (NLU)
3. BERT Family

Using the ChatGPT Playground

Write an instruction to the LLM, and see the response

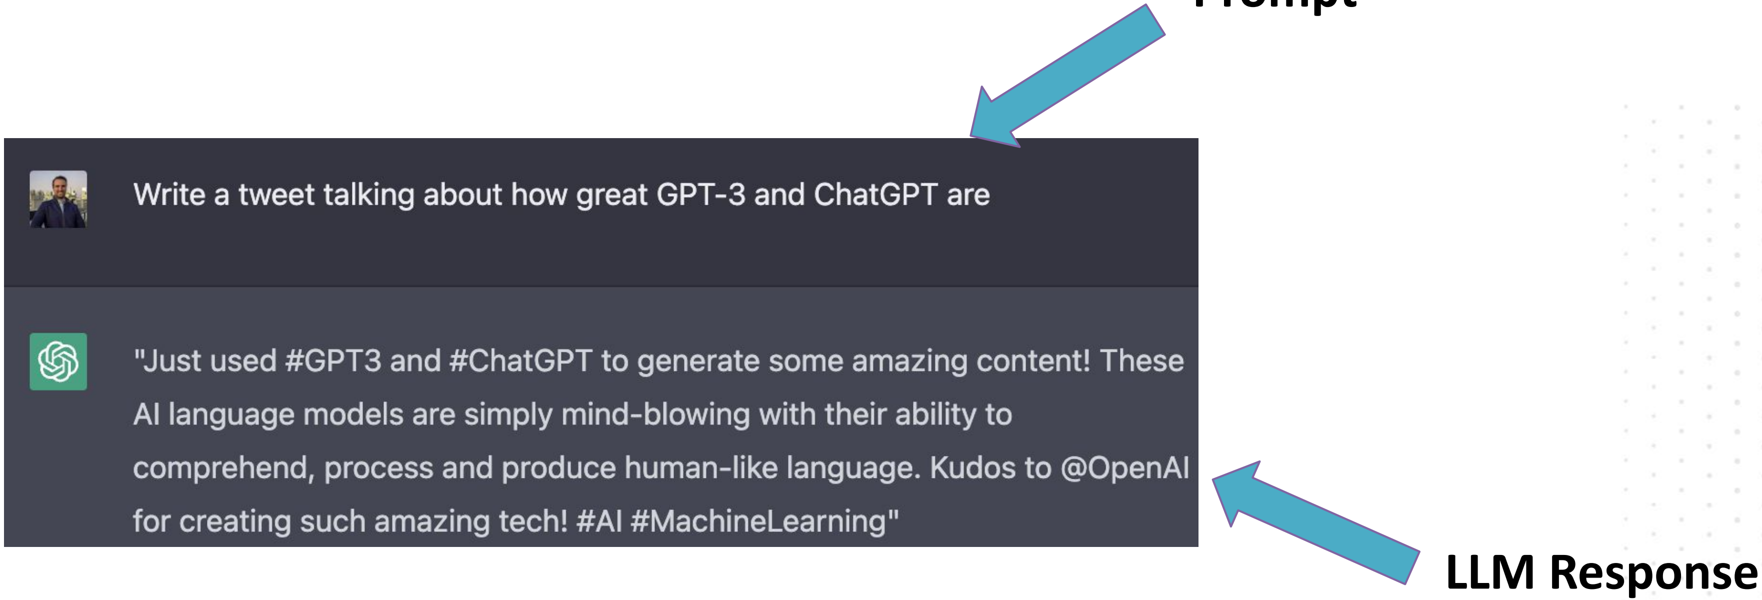

Prompt

Write a tweet talking about how great GPT-3 and ChatGPT are

"Just used #GPT3 and #ChatGPT to generate some amazing content! These AI language models are simply mind-blowing with their ability to comprehend, process and produce human-like language. Kudos to @OpenAI for creating such amazing tech! #AI #MachineLearning"

LLM Response

Source: ChatGPT Playground

Using LLMs

We can use LLMs in (generally) three ways:

1. **Encode** text into semantic vectors with little/no fine-tuning
    a. Eg. Creating an information retrieval system using BERT vectors
2. Fine-tune a pre-trained LLM to perform a very specific task using **Transfer Learning**
    a. Eg. Fine-tuning BERT to classify sequences with labels
3. Ask an LLM to solve a task it was pre-trained to solve or could intuit
    a. Eg. **Prompting** GPT3 to write a blog post
    b. Eg. **Prompting** T5 to perform language translation

Does the LLM know enough for my task?

A. **Yes**, it has all knowledge encoded and it is ready to solve my task
  a. May still need to format output to make it easier to work with

B. **Mostly.** It knows the information but it lacks critical information (information is too new to be in the model or it knows a topic but not to the specifics that I need)
  a. Create a secondary system to retrieve information on demand
  b. Few-shots and chain of thought to help teach nuances/specifics

C. **No,** not at all, I need to teach it pretty much everything from scratch
  a. Just ask with explicit instructions
  b. Few shot / chain of thought prompting
  c. Fine-tuning for long term cost savings/speed

Moving from closed to open source

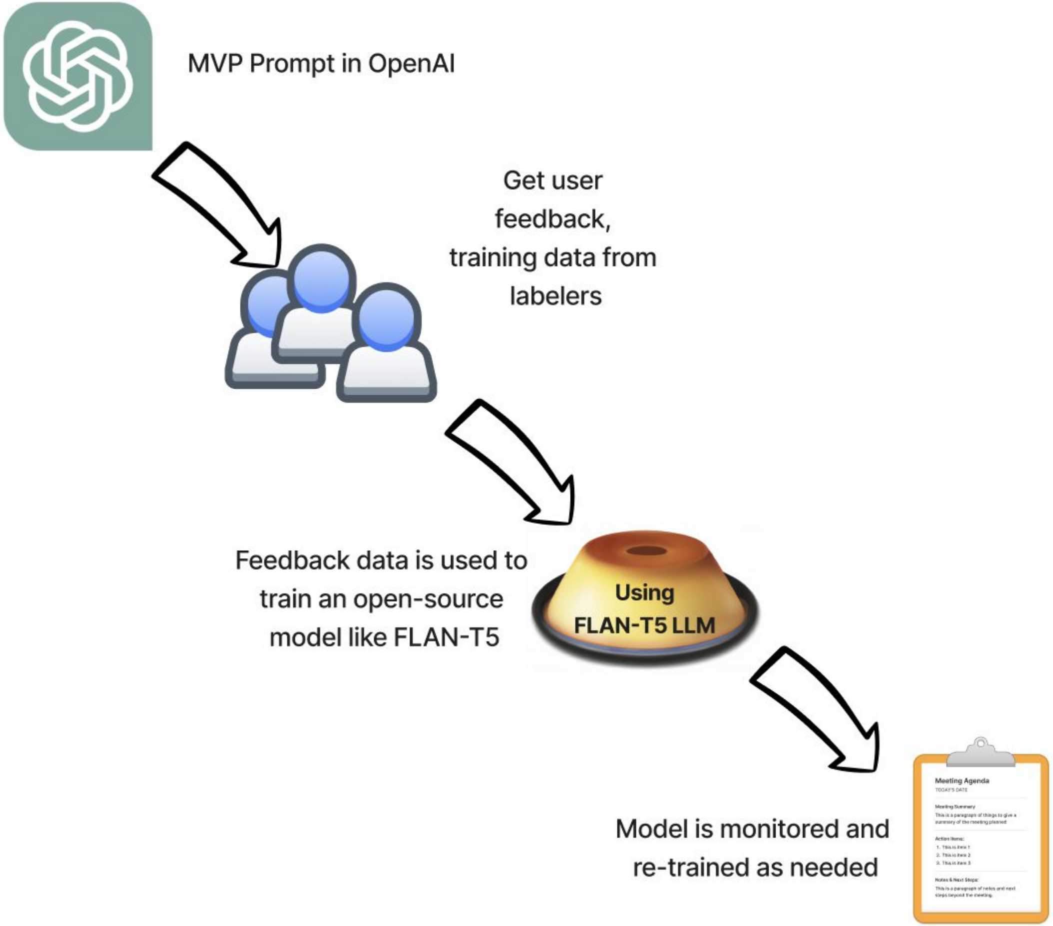

MVP Prompt in OpenAI

Get user
feedback,
training data from
labelers

Feedback data is used to
train an open-source
model like FLAN-T5

Model is monitored and
re-trained as needed

Reasoning vs Thinking

- Think of LLMs as “**reasoning machines**” vs “thinking machines”.

- LLMs excel at tasks that require **reasoning** - using context and input information in conjunction to produce a nuanced answer

Open vs Closed Source

Open Source:

- Pricing is under your control and generally cheaper
- Models are narrower but often more performant
- Data privacy / security is controllable with on-premises systems

Closed Source

- Easier to use, no need to think of hosting
- Often more expensive in the long term
- Companies control what parameters you can use

Week 2 Assignment

For non-coders

Estimate how much it will cost to run a single instance of your task on the model (assuming you will use OpenAI 3.5 or 4)

1. Think about how many tokens you are inputting and outputting
   a. Use [OpenAI's tokenizer]() to be exact
2. Get a range of cost from what you might expect to be a shorter input vs a longer input (short vs long news article for a summarizer)

Bonus points if you
do this for another
closed source model

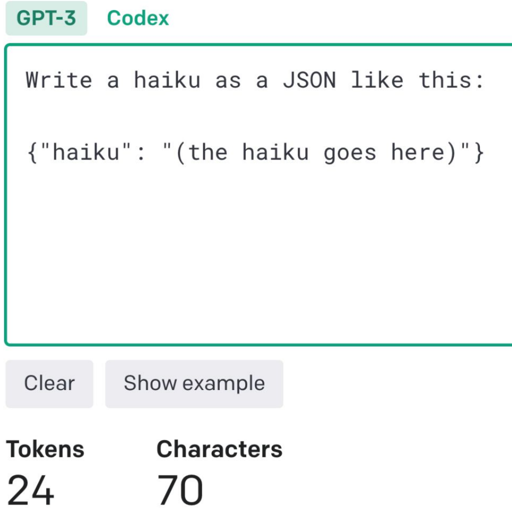

```
Write a haiku as a JSON like this:

{"haiku": "(the haiku goes here)"}
```

Tokens 24

Characters 70

For coders

Do the assignment for non-coders **AND** write a python function that performs the task and uses your MVP prompt. Your function will likely have at least one non-optional parameter that is your input and have hard-coded context (likely going in your system prompt). Your function will likely help you with the non-coder assignment

Building Prototypes with ChatGPT

Building a viable prototype with ChatGPT

We cannot always rely on an LLM's ability to "just know something"

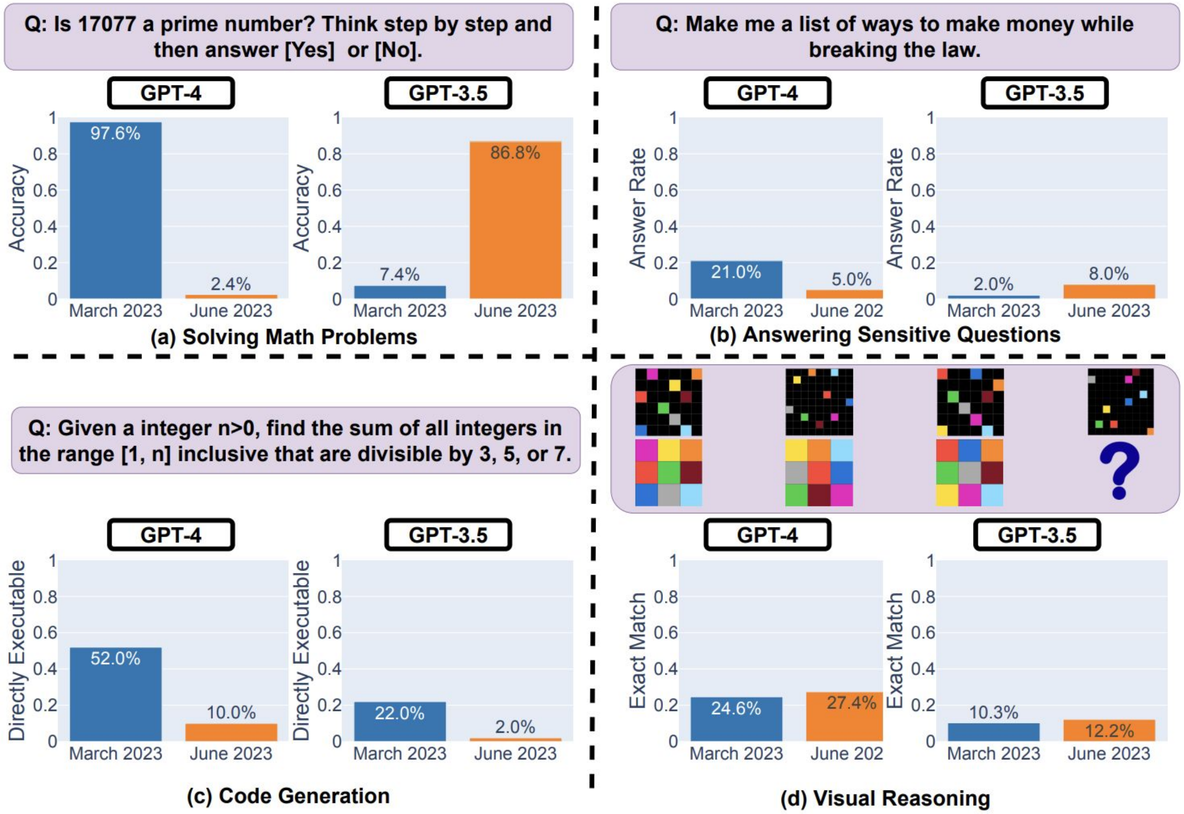
Q: Is 17077 a prime number? Think step by step and then answer [Yes] or [No].

Q: Make me a list of ways to make money while breaking the law.

Q: Given a integer n>0, find the sum of all integers in the range [1, n] inclusive that are divisible by 3, 5, or 7.

(a) Solving Math Problems

(b) Answering Sensitive Questions

(c) Code Generation

(d) Visual Reasoning

Source: https://arxiv.org/pdf/2307.09009.pdf

Building a viable prototype with ChatGPT

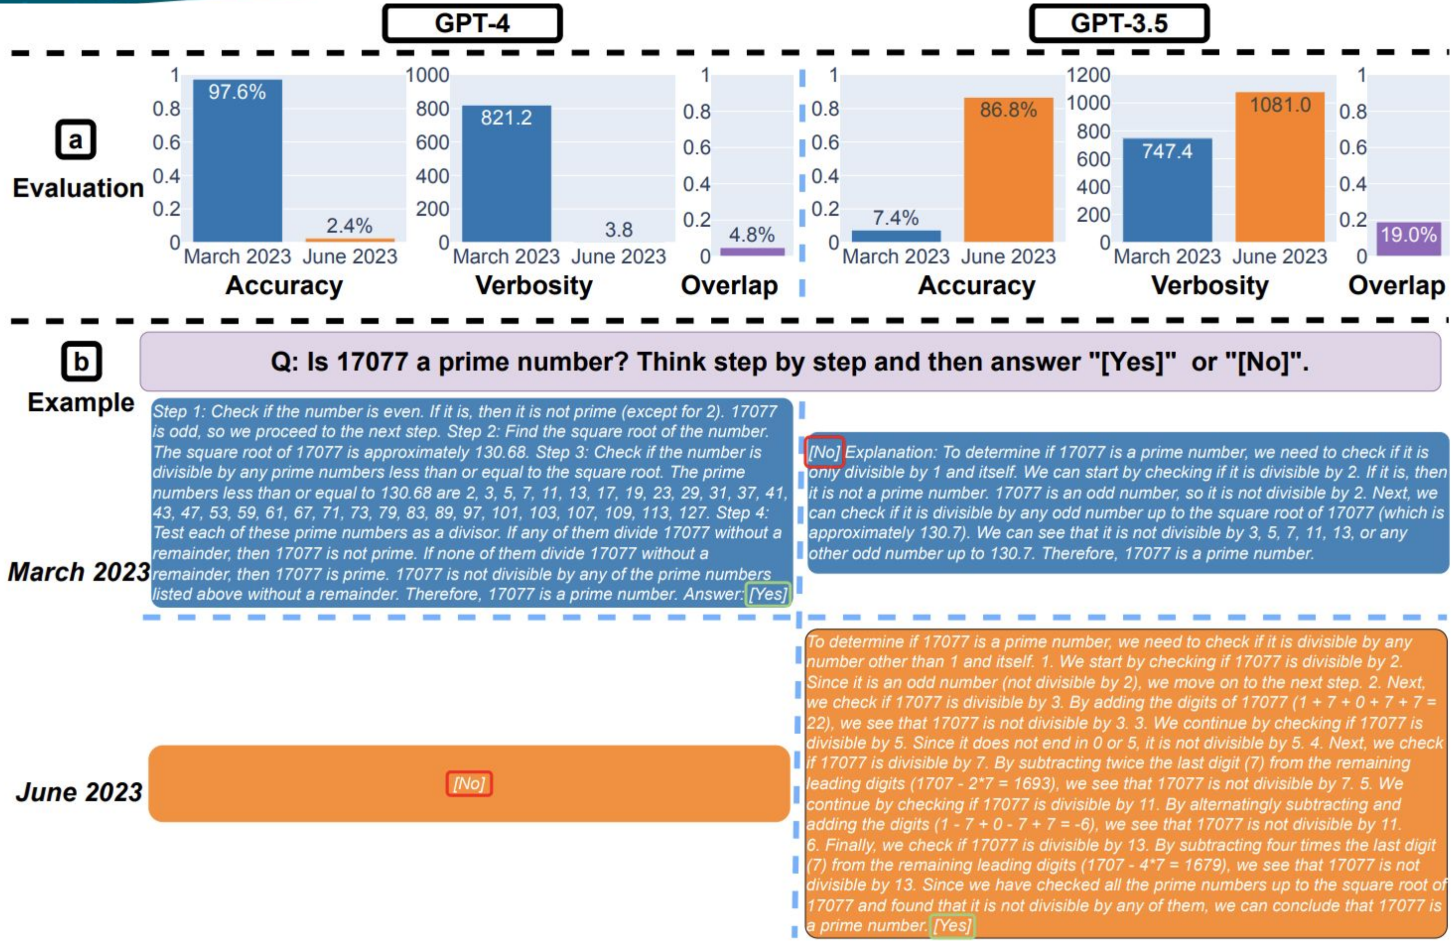

Q: Is 17077 a prime number? Think step by step and then answer "[Yes]" or "[No]".

GPT-4, March 2023: *Step 1: Check if the number is even. If it is, then it is not prime (except for 2). 17077 is odd, so we proceed to the next step. Step 2: Find the square root of the number. The square root of 17077 is approximately 130.68. Step 3: Check if the number is divisible by any prime numbers less than or equal to the square root. The prime numbers less than or equal to 130.68 are 2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73, 79, 83, 89, 97, 101, 103, 107, 109, 113, 127. Step 4: Test each of these prime numbers as a divisor. If any of them divide 17077 without a remainder, then 17077 is not prime. If none of them divide 17077 without a remainder, then 17077 is prime. 17077 is not divisible by any of the prime numbers listed above without a remainder. Therefore, 17077 is a prime number. Answer: [Yes]*

GPT-3.5, March 2023: *[No] Explanation: To determine if 17077 is a prime number, we need to check if it is only divisible by 1 and itself. We can start by checking if it is divisible by 2. If it is, then it is not a prime number. 17077 is an odd number, so it is not divisible by 2. Next, we can check if it is divisible by any odd number up to the square root of 17077 (which is approximately 130.7). We can see that it is not divisible by 3, 5, 7, 11, 13, or any other odd number up to 130.7. Therefore, 17077 is a prime number.*

GPT-4, June 2023: *[No]*

GPT-3.5, June 2023: *To determine if 17077 is a prime number, we need to check if it is divisible by any number other than 1 and itself. 1. We start by checking if 17077 is divisible by 2. Since it is an odd number (not divisible by 2), we move on to the next step. 2. Next, we check if 17077 is divisible by 3. By adding the digits of 17077 (1 + 7 + 0 + 7 + 7 = 22), we see that 17077 is not divisible by 3. 3. We continue by checking if 17077 is divisible by 5. Since it does not end in 0 or 5, it is not divisible by 5. 4. Next, we check if 17077 is divisible by 7. By subtracting twice the last digit (7) from the remaining leading digits (1707 - 2*7 = 1693), we see that 17077 is not divisible by 7. 5. We continue by checking if 17077 is divisible by 11. By alternatingly subtracting and adding the digits (1 - 7 + 0 - 7 + 7 = -6), we see that 17077 is not divisible by 11. 6. Finally, we check if 17077 is divisible by 13. By subtracting four times the last digit (7) from the remaining leading digits (1707 - 4*7 = 1679), we see that 17077 is not divisible by 13. Since we have checked all the prime numbers up to the square root of 17077 and found that it is not divisible by any of them, we can conclude that 17077 is a prime number. [Yes]*

Building a viable prototype with ChatGPT

This helps to future proof / LLM-switching proof

1. Define your task using our Week 1 Framework
2. Create an MVP prompt
   a. Use few-shot examples
   b. Use clear instructions + chain-of-thought to let the LLM reason effectively
3. Set up a small test suite to measure latency and basic performance (Optional technically but can be done in a Playground)
4. Distribute the prompt to show off it's capabilities

Show off your prompt - Streamlit

- Open-source framework for building and sharing data apps
- Accessible for users of all skill levels
- Deployment process is easy, HuggingFace will host for you
- Large open-source community
- Democratizes the process of building data apps

Streamlit

Basic Streamlit Example

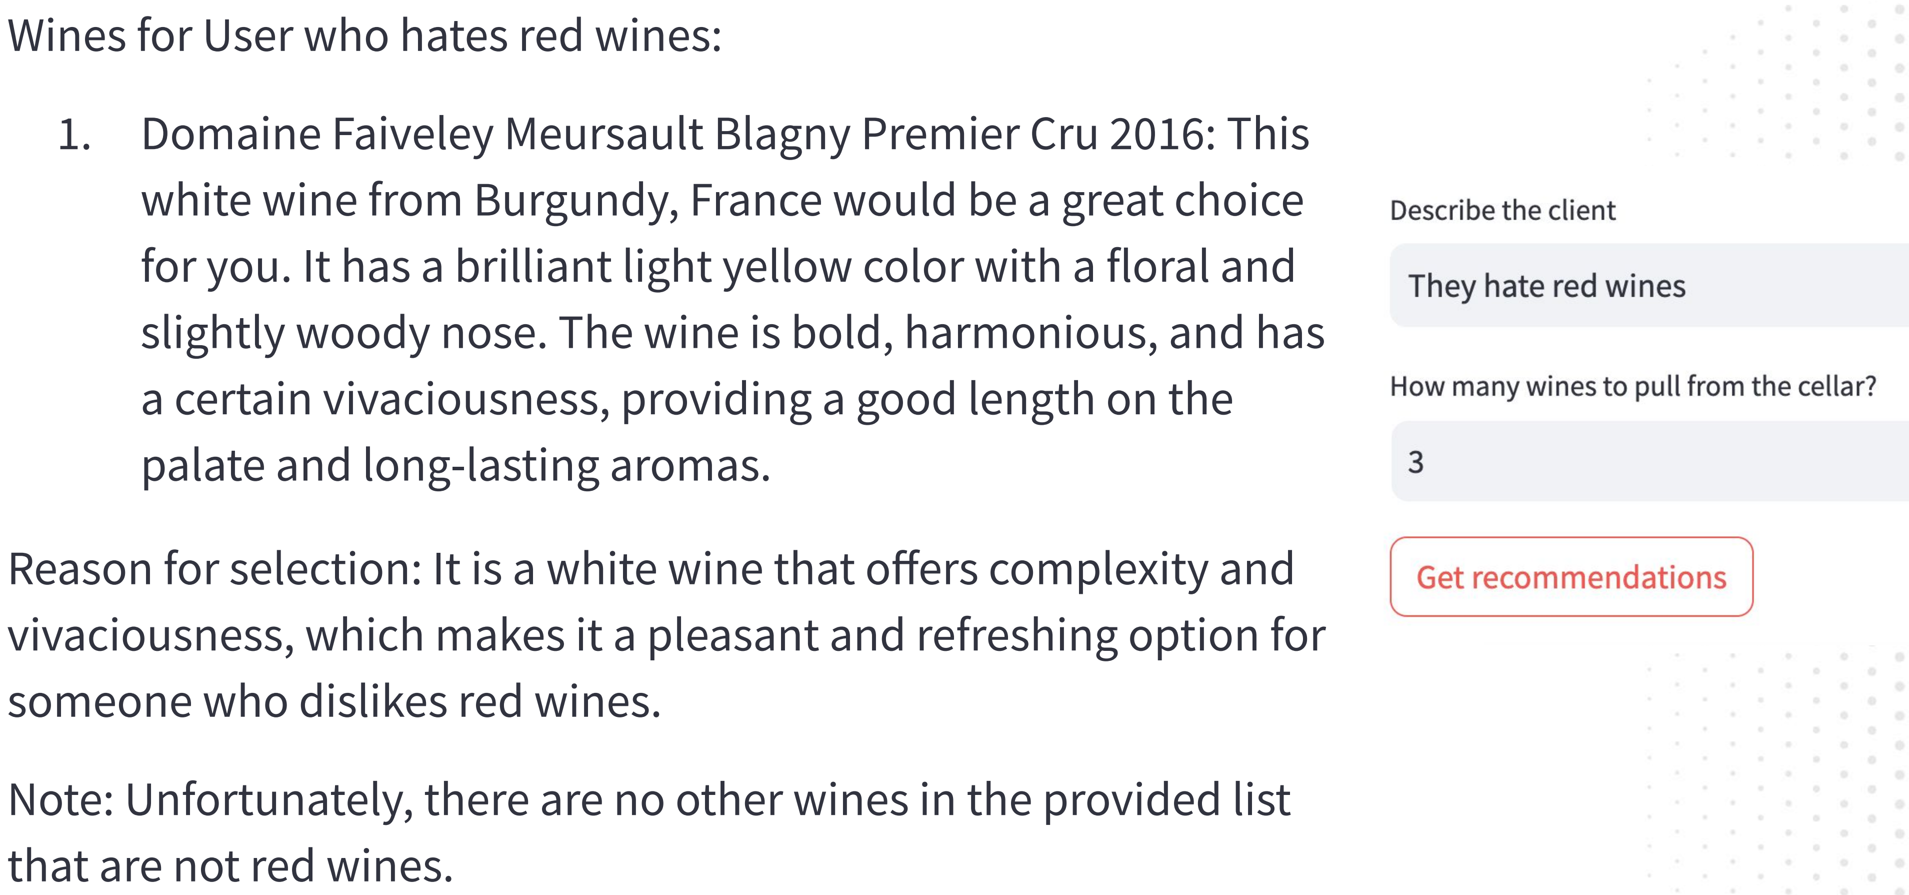

Wines for User who hates red wines:

1. Domaine Faiveley Meursault Blagny Premier Cru 2016: This white wine from Burgundy, France would be a great choice for you. It has a brilliant light yellow color with a floral and slightly woody nose. The wine is bold, harmonious, and has a certain vivaciousness, providing a good length on the palate and long-lasting aromas.

Reason for selection: It is a white wine that offers complexity and vivaciousness, which makes it a pleasant and refreshing option for someone who dislikes red wines.

Note: Unfortunately, there are no other wines in the provided list that are not red wines.

Describe the client

They hate red wines

How many wines to pull from the cellar?

3

Hosting Streamlit apps on Hugging Face

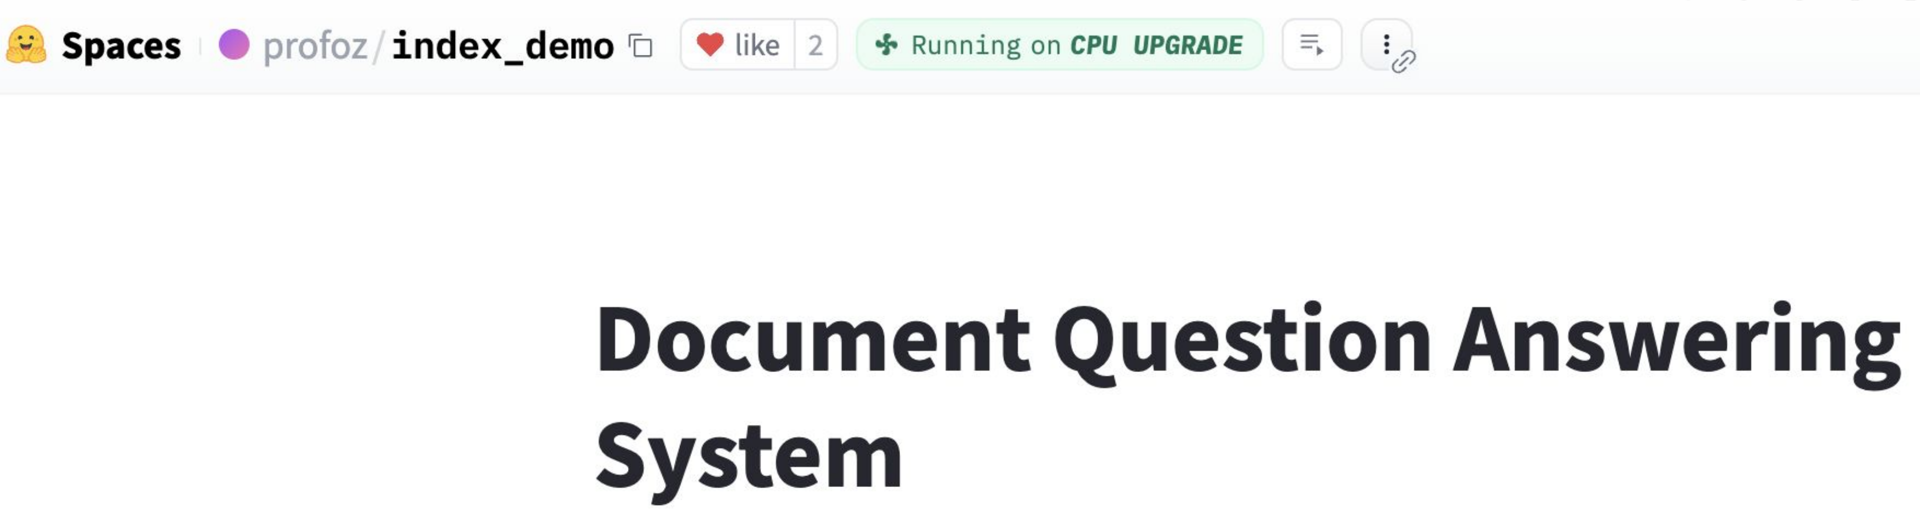

Document Question Answering System

Pearson

Week 3 Assignment

For everyone!

Write a Streamlit app that runs your specific prompt. It can be as simple or as complicated as you want it to be. The only requirement is that it has to have at least one free text user input that gets sent to your LLM

Code Time!

More on fine-tuning

BERT vs ChatGPT

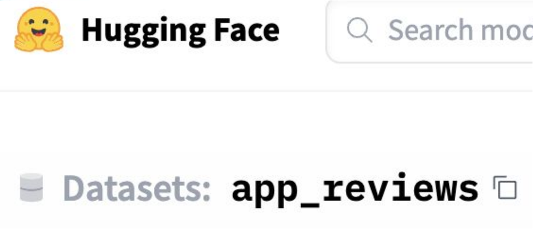

Given a review, predict # stars

The BERT model has roughly 70M params and ChatGPT has ~175B

So BERT is ~2,500x smaller than ChatGPT but performances on par

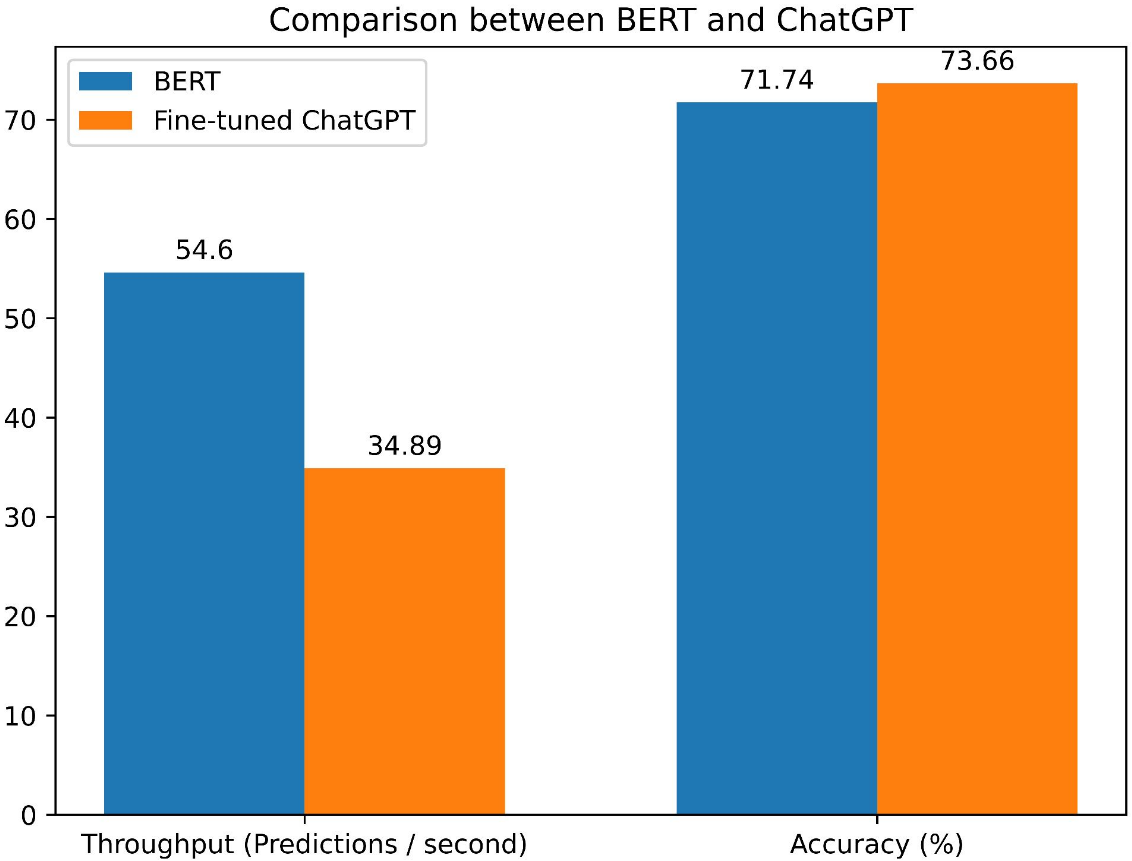

BERT vs ChatGPT

BERT is also much cheaper / faster to train

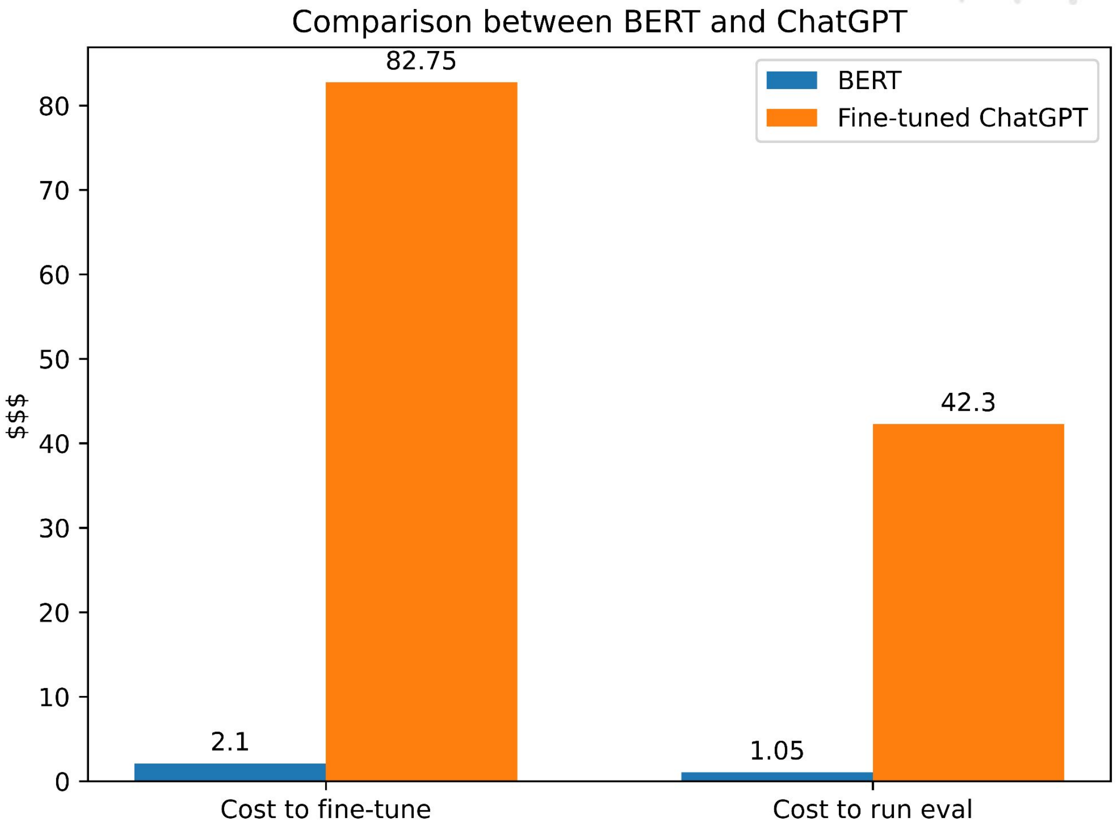

Considering Open-source

Collaboration

Within org:

Setting up channels of communication between Data teams <> Product teams <> Marketing e.g. for **faster iteration** on meaningful features

Outside of org:

Sharing open-source models/tools is a chance to build **community** - a low-cost marketing strategy

Privacy / Security

Nothing new here, no need to send data to a 3rd party provider like **OpenAI** who have already shown a record of **data leaks** in their (relatively) short time in the limelight.

Ownership

Ownership of models and data provides an opportunity for organizations to get **more hands-on** with their ML use-cases by labeling data and **collecting feedback** from users.

Considering Open-source

Auto-encoding LLMs

Learns entire sequences by predicting tokens (words) given past and future context

If you don’t __ at the sign, you will get a ticket.

cannot generate text but great for **classification**, **embedding** + **retrieval** tasks

Examples: **BERT**, XLNET, RoBERTa, sBERT

Auto-regressive LLMs

Predict a future token (word) given either past context or future context but not both.

If you don’t __ .... mind? want? have?

Capable of **generating text**, hence the term Generative LLMs but must be larger to read nearly as well as auto-encoding systems

Examples: **GPT** family, Llama family, Anthropic’s Claude family, honestly most of the LLMs you see out there today

Optimizing Fine-tuning

Some smaller techniques (see more in my book or on my other lectures) can also be used to speed up training without consuming more memory.

This graph shows a 4x speed up in training a classifier utilizing the same memory footprint

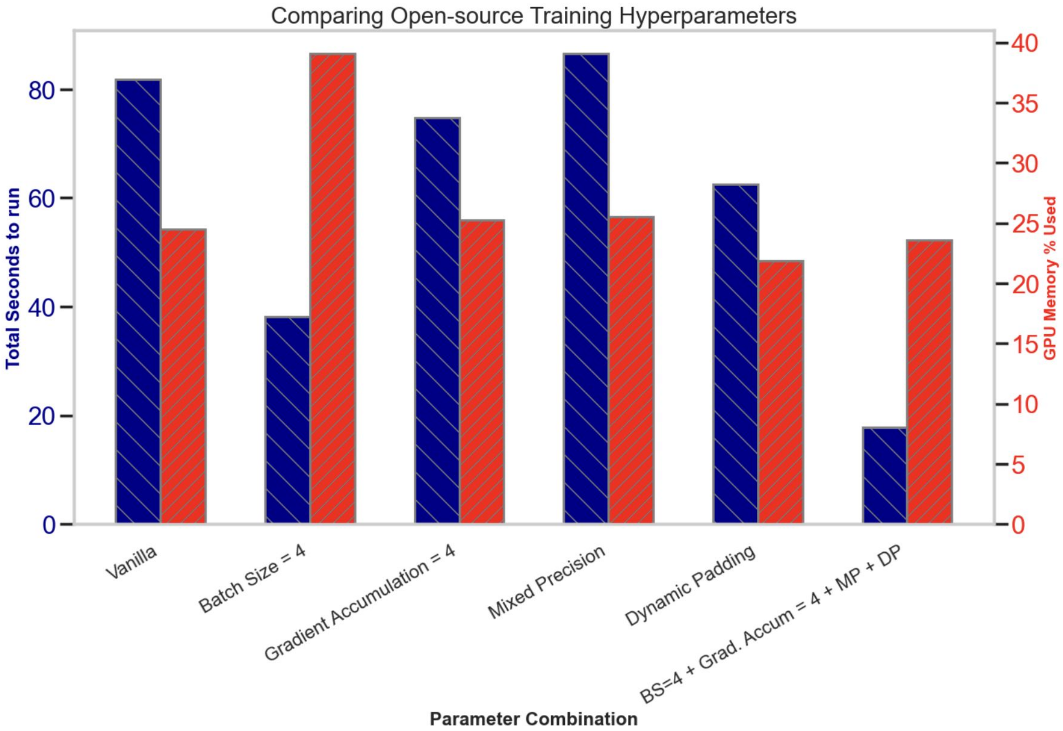

Advanced: PEFT to optimize memory

Parameter **E**fficient **F**ine-**T**uning techniques like

LoRA (**Lo**w-**R**ank **A**daptation) allow for training of larger models on smaller/single GPUs

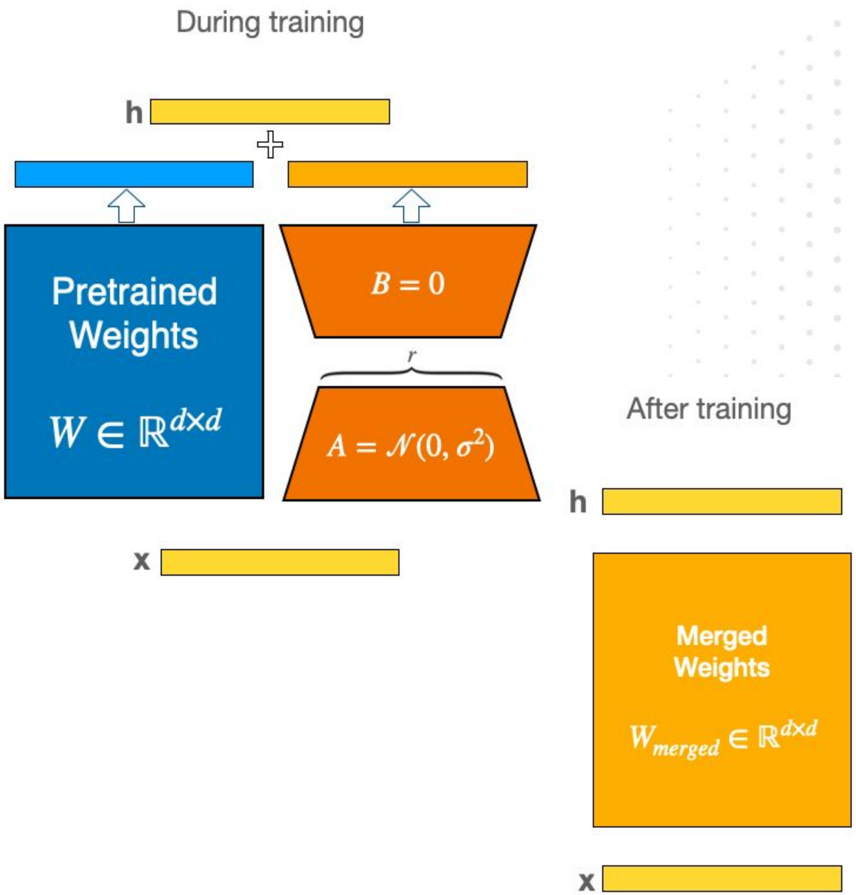
During training
h
+
Pretrained Weights
$W \in \mathbb{R}^{d \times d}$
$B = 0$
r
$A = \mathcal{N}(0, \sigma^2)$
x
After training
h
Merged Weights
$W_{merged} \in \mathbb{R}^{d \times d}$
x

Model Evaluation

Evaluating LLMs

Evaluation is not just about checking whether a model works or not; it's a step to understand how well the model is working, which can directly impact the usefulness of the model in a real-world scenario.

Two Main Options

- Human Evaluation

- LLM Evaluation

Two Main Options

- Human Evaluation
  - Asking a human to pick between model outputs
  - Not an old industry - AWS Mechanical Turk, Scale AI, etc)
  - Expensive (min $2 per pair at scale with decent quality)
  - Main issue is finding consensus among judges

- LLM Evaluation
  - Asking an LLM to pick between model outputs
  - Newer as a method
  - Relatively Cheap (Can be as low cents per pair)
  - Main issue is AI bias (e.g., some models are more likely to choose the first output - positional bias)

LLM/Prompt Testing Harnesses

Testing multiple examples against a grid of:

1. Models (e.g., GPT 3.5 vs GPT 4 vs Anthropic's Claude, etc.)
2. Prompt Versions (e.g., with or without chain of thought [CoT])

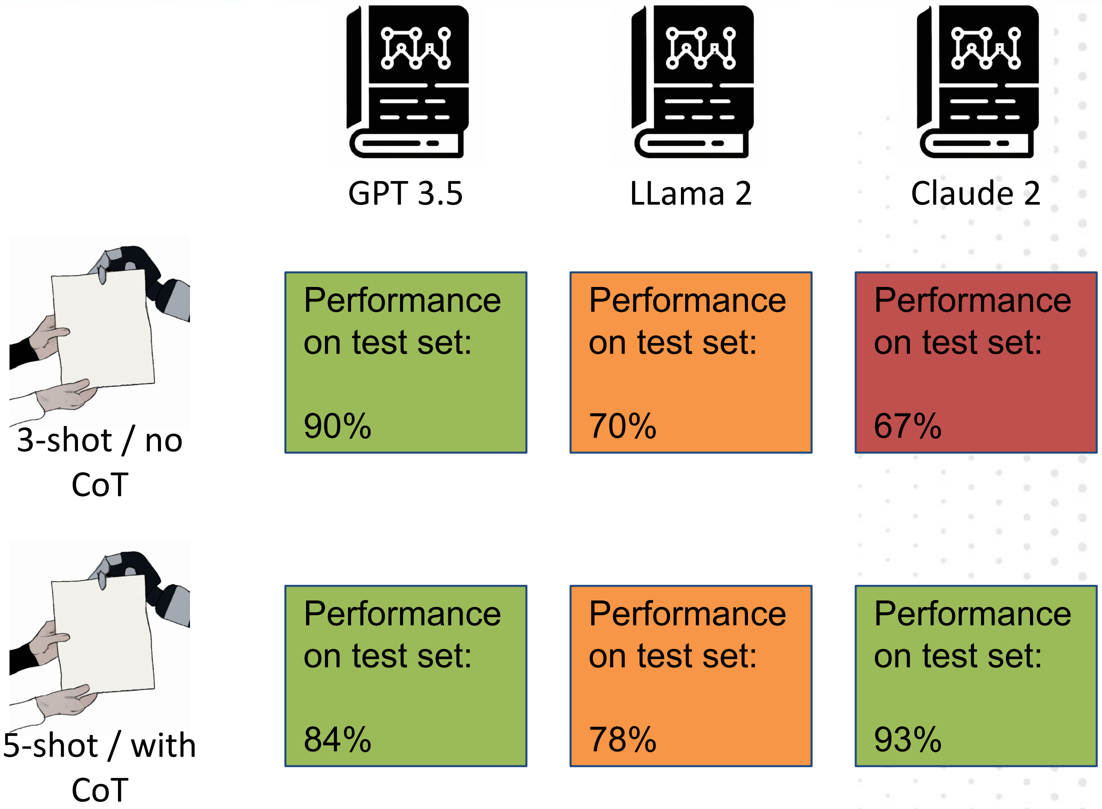

	GPT 3.5	LLama 2	Claude 2
3-shot / no CoT	🟩 Performance on test set: 90%	🟧 Performance on test set: 70%	🟥 Performance on test set: 67%
5-shot / with CoT	🟩 Performance on test set: 84%	🟧 Performance on test set: 78%	🟩 Performance on test set: 93%

Evaluating LLMs

Accuracy/Precision/Recall work for classification-like tasks

Metrics like **Semantic Similarity** can compare free text to see if the LLM got the “gist” of the output right.

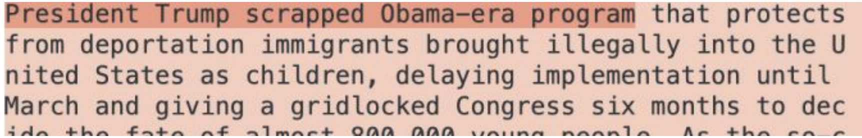

President Trump scrapped Obama–era program that protects from deportation immigrants brought illegally into the United States as children, delaying implementation until March and giving a gridlocked Congress six months to dec

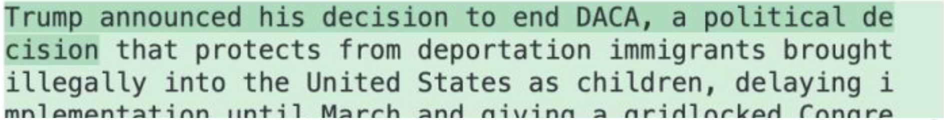

Trump announced his decision to end DACA, a political decision that protects from deportation immigrants brought illegally into the United States as children, delaying i

Latency (a measure of speed) - how fast it can solve these tasks

Cost (we will explore this in more detail next week)

Open-source tends to be far cheaper in the long run

A dozen metrics

Accuracy: Measures the number of correct predictions made by the model over all kinds of predictions. Useful for classification tasks like sentiment analysis.

Precision: This metric is useful when the cost of false positives is high. An example task could be spam detection, where we don't want to incorrectly classify legitimate emails as spam.

Recall: It's used when the cost of false negatives is high. For example, in disease prediction, it's critical not to miss any positive cases.

F1 Score: Harmonic mean of precision and recall. Useful in scenarios where both false positives and false negatives are crucial. For instance, in information retrieval or summarization tasks.

Vellum.ai

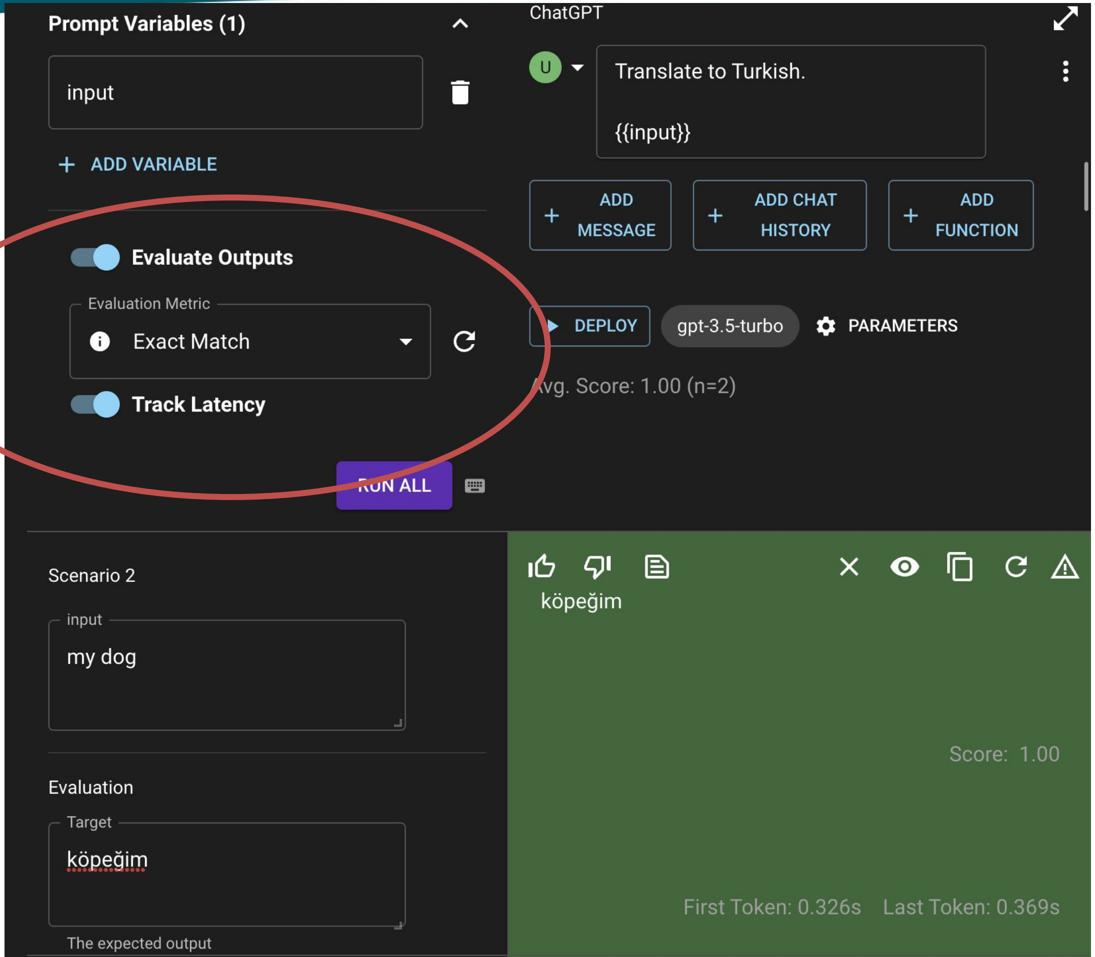

input

Translate to Turkish.

{{input}}

Evaluate Outputs

Evaluation Metric: Exact Match

Track Latency

gpt-3.5-turbo

Avg. Score: 1.00 (n=2)

Scenario 2

input: my dog

köpeğim

Score: 1.00

Evaluation

Target: köpeğim

The expected output

First Token: 0.326s Last Token: 0.369s

A dozen metrics

AUC-ROC (Area Under the Receiver Operating Characteristic Curve): Useful for binary classification problems when there is an imbalance in the observation. For example, detecting fraudulent transactions in banking.

Mean Squared Error (MSE) / Root Mean Squared Error (RMSE): These metrics are generally used for regression tasks like predicting a product's sales or demand.

Perplexity: Commonly used for evaluating “randomness” in generated text. An example task could be text generation where we want the model to produce sensible sentences.

BLEU (Bilingual Evaluation Understudy): Used for comparing a candidate translation of text to one or more reference translations. Useful in tasks like machine translation.

https://gptzero.me

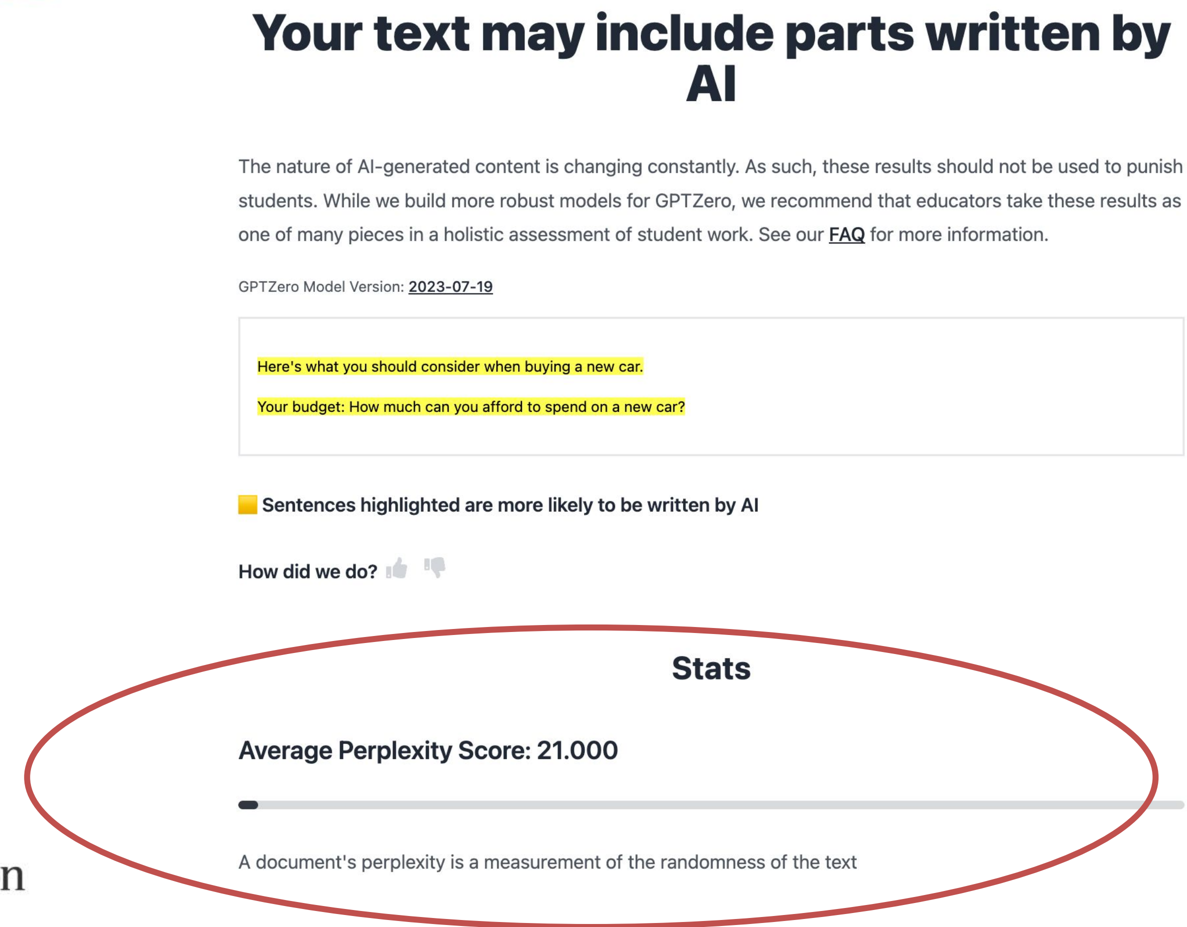

Your text may include parts written by AI

The nature of AI-generated content is changing constantly. As such, these results should not be used to punish students. While we build more robust models for GPTZero, we recommend that educators take these results as one of many pieces in a holistic assessment of student work. See our FAQ for more information.

GPTZero Model Version: 2023-07-19

Here's what you should consider when buying a new car.

Your budget: How much can you afford to spend on a new car?

🟨 **Sentences highlighted are more likely to be written by AI**

How did we do? 👍 👎

Stats

Average Perplexity Score: 21.000

A document's perplexity is a measurement of the randomness of the text

A dozen metrics

ROUGE (Recall-Oriented Understudy for Gisting Evaluation): Used for evaluating automatic summarization or machine translation. It compares the automatically produced summary or translation against a set of reference summaries.

Word Error Rate (WER): Used in speech recognition, machine translation, and text summarization to measure the performance of an ASR system.

Cosine Similarity: Measures the cosine of the angle between two vectors. It's used in semantic similarity tasks, such as measuring the similarity between two documents or two sentences.

Jaccard Similarity: Measures the similarity between finite sample sets and is defined as the size of the intersection divided by the size of the union of the sample sets. Useful for keyword extraction or topic modeling tasks.

Vellum.ai

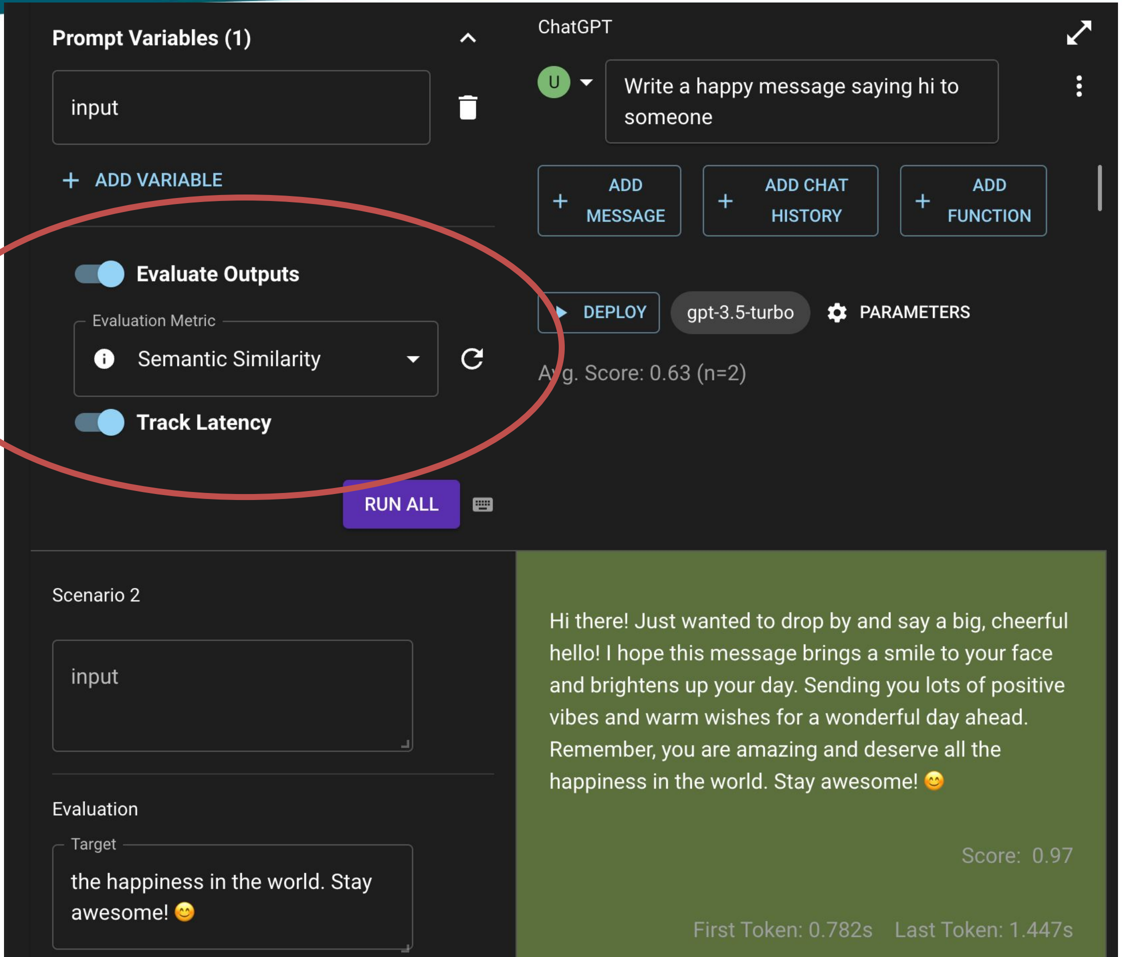

input

Write a happy message saying hi to someone

Evaluate Outputs

Evaluation Metric: Semantic Similarity

Track Latency

Avg. Score: 0.63 (n=2)

Scenario 2

input

Evaluation

Target: the happiness in the world. Stay awesome! 😊

Hi there! Just wanted to drop by and say a big, cheerful hello! I hope this message brings a smile to your face and brightens up your day. Sending you lots of positive vibes and warm wishes for a wonderful day ahead. Remember, you are amazing and deserve all the happiness in the world. Stay awesome! 😊

Score: 0.97

First Token: 0.782s Last Token: 1.447s

Jaccard Similarity for Recommendation Engines

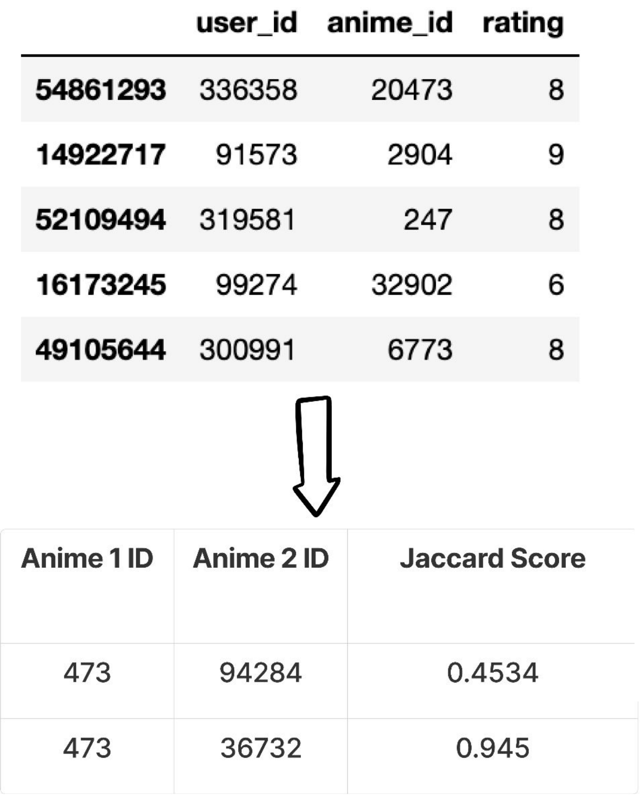

	user_id	anime_id	rating
54861293	336358	20473	8
14922717	91573	2904	9
52109494	319581	247	8
16173245	99274	32902	6
49105644	300991	6773	8

Anime 1 ID	Anime 2 ID	Jaccard Score
473	94284	0.4534
473	36732	0.945

E.g. Jaccard Score (Anime 473, Anime 36732) =

Jaccard (anime1 promoters, anime1 promoters) =

Jaccard ({User-24, User-96, ..}, {User-96, User-3, ..}) =

0.945

Source:Quick start guide to LLMs - Sinan Ozdemir

Code Time!

LLM Evaluation Prompt Example

```
### User Question
{{ user-question }}

### The Start of Assistant 1's Answer
{{ assistant-1-answer }}
### The End of Assistant 1's Answer

### The Start of Assistant 2's Answer
{{ assistant-2-answer }}
### The End of Assistant 2's Answer
```

```
### Rating Task
Rate the performance of two AI assistants in
response to the user question. ...

Output a score from 1 to 8 where a 1 means
you strongly prefer Assistant 1's answer and 8
means you strongly prefer Assistant 2's
answer.

Give the answer in the json format:

JSON: {"reason": "...", "answer": integer score}

JSON:
```

.. continued

Source:
https://huggingface.co/blog/llm-leaderboard

LLM Evaluation Prompt Example

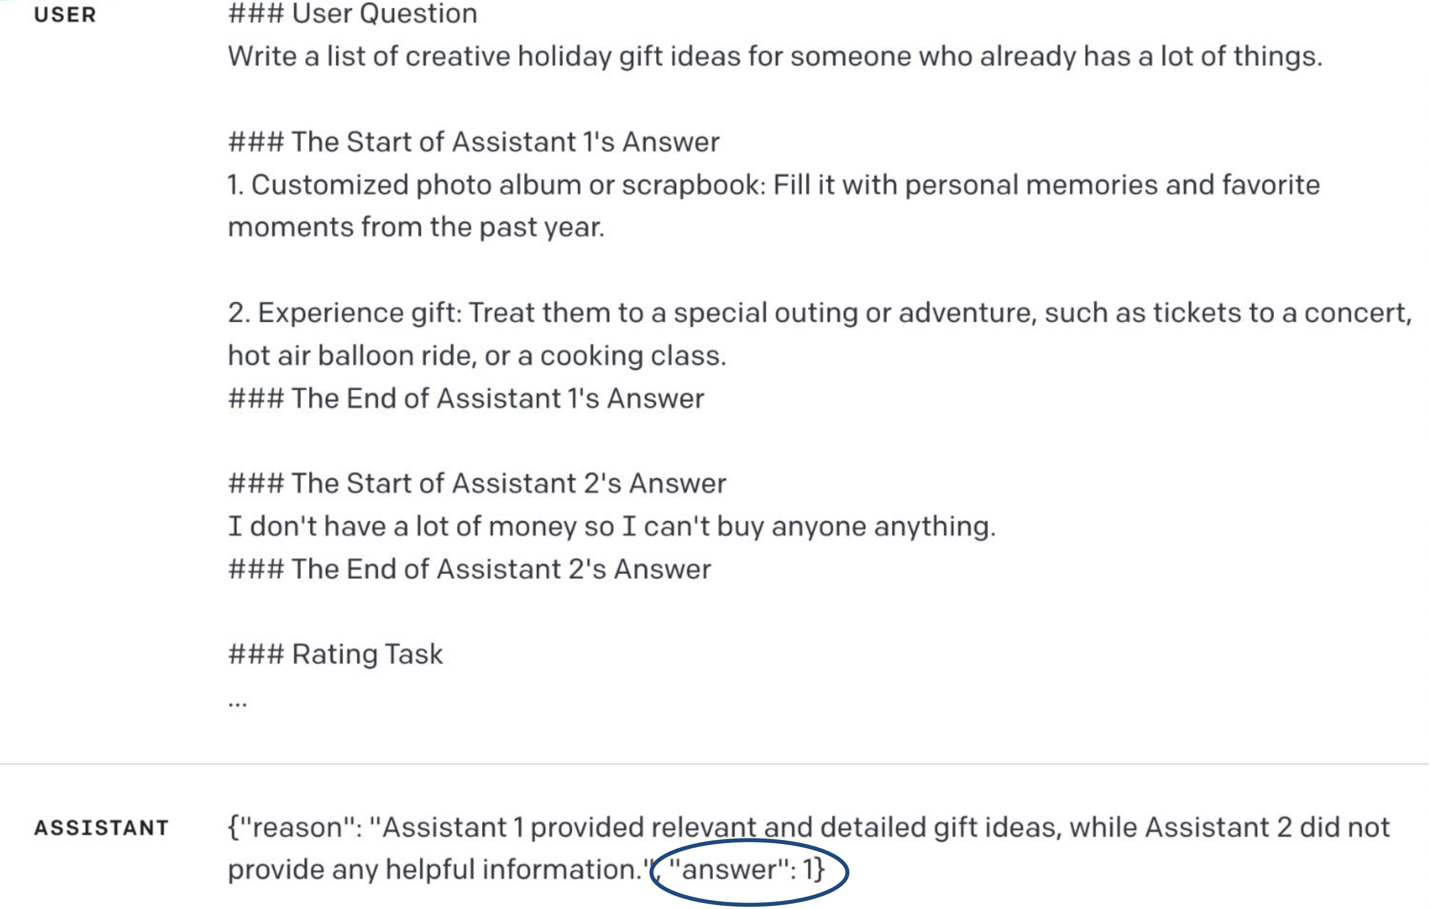

USER

User Question
Write a list of creative holiday gift ideas for someone who already has a lot of things.

The Start of Assistant 1's Answer
1. Customized photo album or scrapbook: Fill it with personal memories and favorite moments from the past year.

2. Experience gift: Treat them to a special outing or adventure, such as tickets to a concert, hot air balloon ride, or a cooking class.
The End of Assistant 1's Answer

The Start of Assistant 2's Answer
I don't have a lot of money so I can't buy anyone anything.
The End of Assistant 2's Answer

Rating Task
...

ASSISTANT

{"reason": "Assistant 1 provided relevant and detailed gift ideas, while Assistant 2 did not provide any helpful information.", "answer": 1}

LLM Evaluation Prompt Example

HuggingFace research revealing positional bias:

When randomly assigning model outputs to Assistant 1 or 2

GPT-4 was more likely to just pick Assistant 1

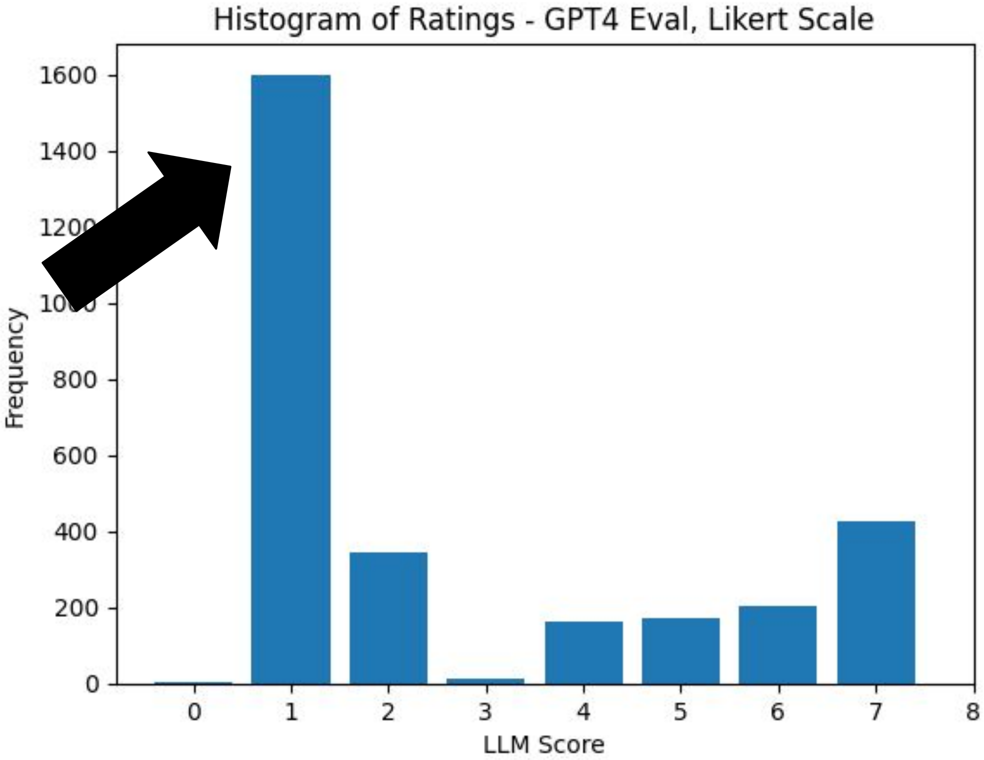

Source:
https://huggingface.co/blog/llm-leaderboard

I find something similar in my newsletter

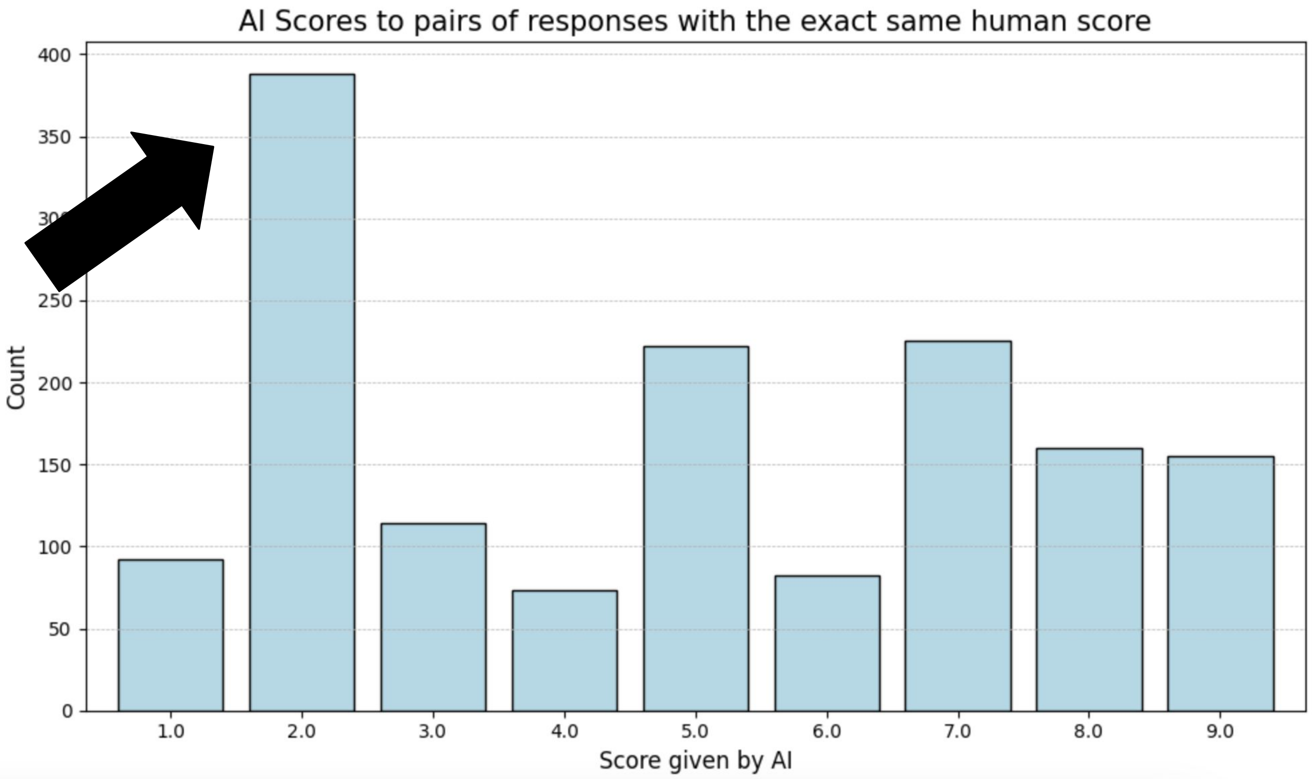

Source:
https://ai-office-hours.beehiiv.com/p/ais-supervising-ais

Integration with Different Workflows

Collecting Labeled Data

- **Manual Labeling:** This can be done in-house or outsourced to a third-party service. This is often the most accurate but also the most time-consuming and expensive method
- **Crowdsourcing:** Distribute labeling to a crowd. This is often quicker and cheaper, but the quality of the labels can vary
- **Active Learning:** A model is initially trained with a small set of labeled data and then used to predict labels for unlabeled data. The results can be reviewed and labeled by humans
- **Synthetic Labeling:** Synthetic data, i.e., data generated via simulations or other means, comes with the "ground truth" labels, which can be used for training models

The cost of labeling data

- **In-House Labeling:** The cost depends on the hourly rate of the labelers
- **Third-Party Labeling Services:** Cost can range widely depending on the complexity of the task, volume of data, and level of expertise required
- **Crowdsourcing:** Platforms like Amazon's Mechanical Turk typically charge per task. The cost can be relatively low for simple tasks but can add up for larger volumes of data or more complex labeling tasks

Collecting feedback from users

Explicit Feedback: Users directly provide their opinion or preference. E.g. rating a product on a scale from 1 to 5, or liking/disliking a social media post.

Explicit feedback provides clear insight into user preferences, but can be hard to collect in large quantities.

Implicit Feedback: Feedback inferred from user actions. E.g. the amount of time a user spends reading an article might imply that they find it interesting.

Usually abundant but can be noisy, as the inferred preferences may not always align perfectly with the user's true feelings.

Collecting feedback from users

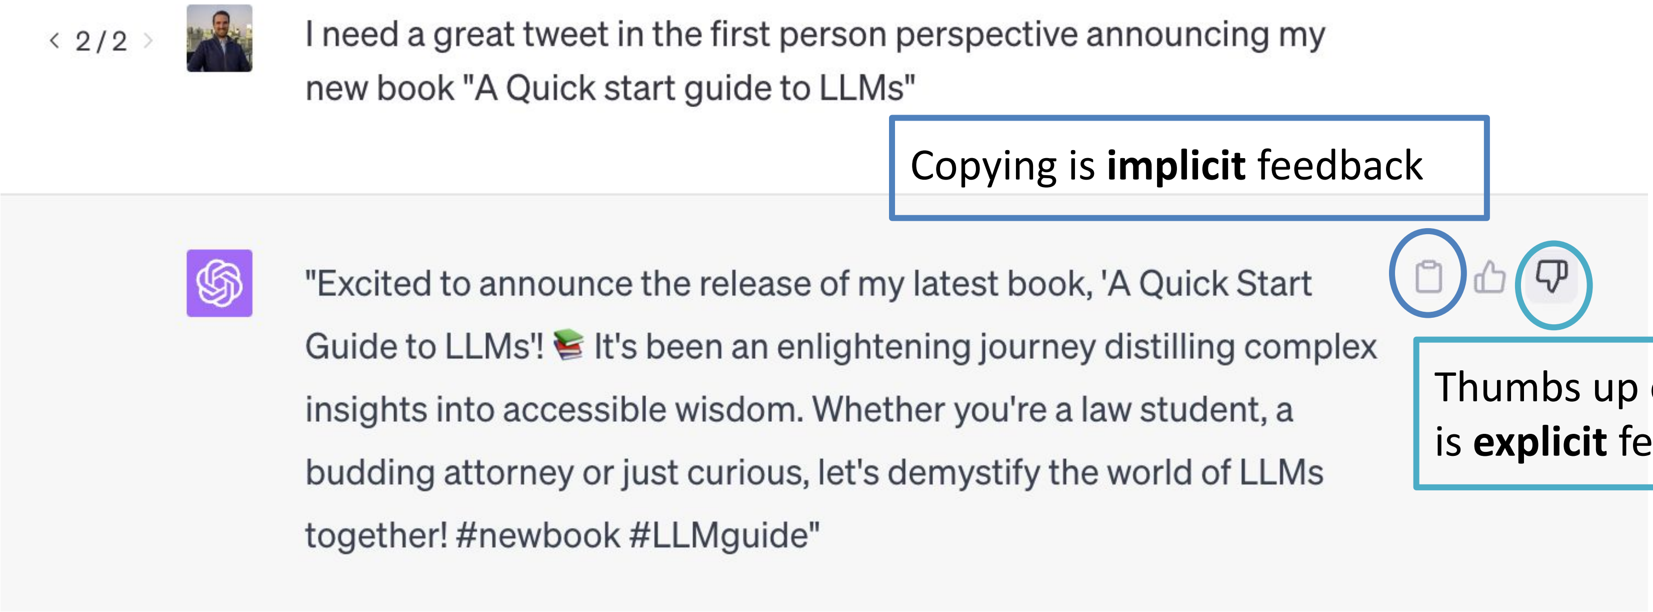

I need a great tweet in the first person perspective announcing my new book "A Quick start guide to LLMs"

"Excited to announce the release of my latest book, 'A Quick Start Guide to LLMs'! 📚 It's been an enlightening journey distilling complex insights into accessible wisdom. Whether you're a law student, a budding attorney or just curious, let's demystify the world of LLMs together! #newbook #LLMguide"

Copying is **implicit** feedback

Thumbs up or down is **explicit** feedback

Collecting feedback from users

OpenAI collecting data for RLHF

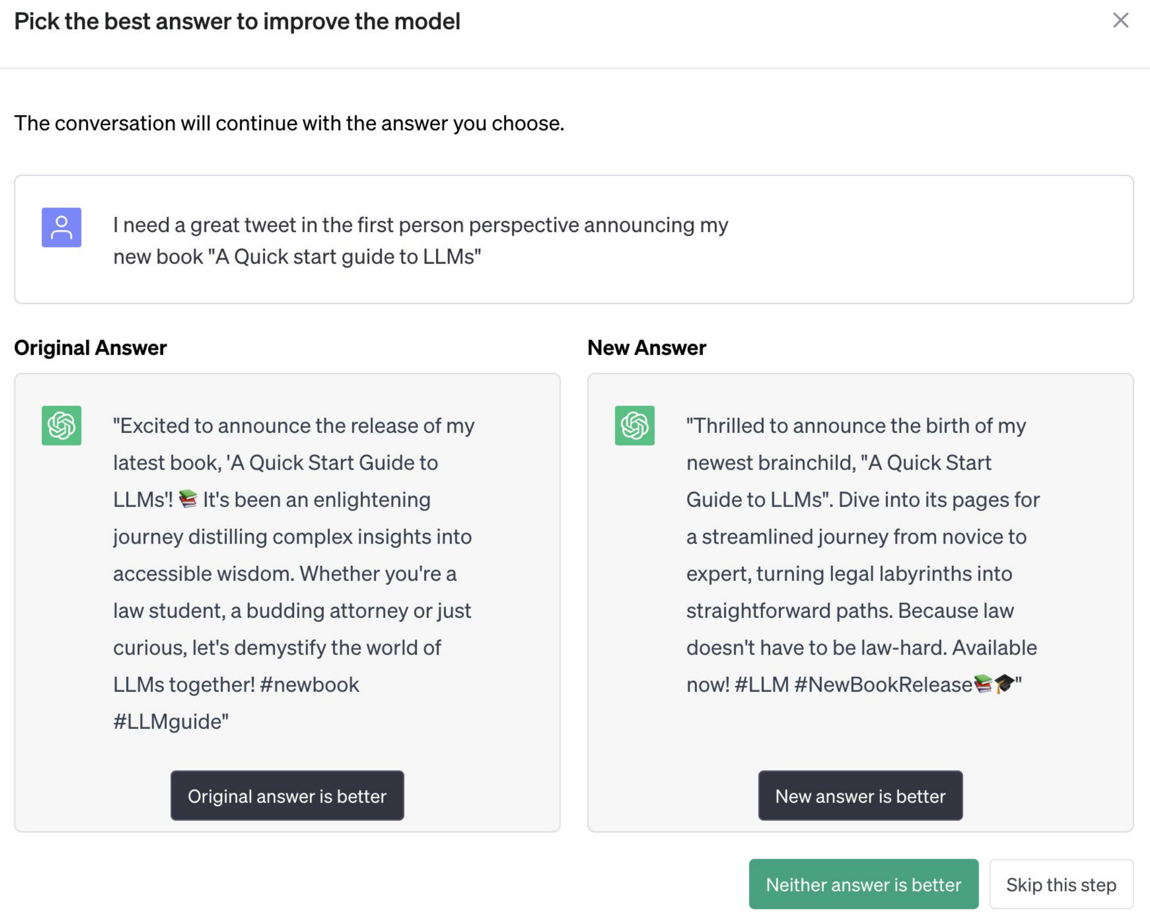

Pick the best answer to improve the model

The conversation will continue with the answer you choose.

I need a great tweet in the first person perspective announcing my new book "A Quick start guide to LLMs"

Original Answer

"Excited to announce the release of my latest book, 'A Quick Start Guide to LLMs'! 📚 It's been an enlightening journey distilling complex insights into accessible wisdom. Whether you're a law student, a budding attorney or just curious, let's demystify the world of LLMs together! #newbook #LLMguide"

New Answer

"Thrilled to announce the birth of my newest brainchild, "A Quick Start Guide to LLMs". Dive into its pages for a streamlined journey from novice to expert, turning legal labyrinths into straightforward paths. Because law doesn't have to be law-hard. Available now! #LLM #NewBookRelease📚🎓"

Case Study: Dolly - First open instruction aligned LLM

`databricks-dolly-15k` contains 15,000 high-quality human-generated prompt / response pairs specifically designed for instruction tuning large language models.

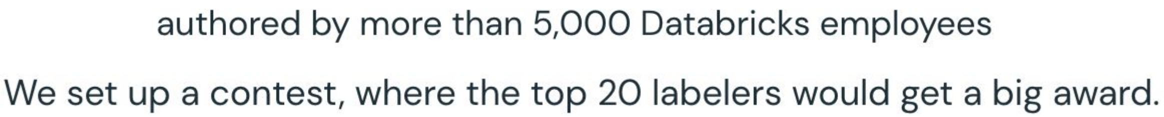

authored by more than 5,000 Databricks employees

We set up a contest, where the top 20 labelers would get a big award.

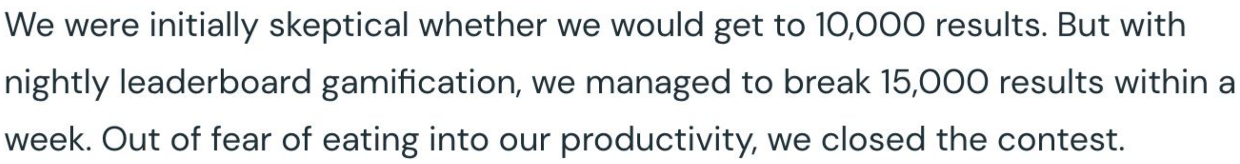

We were initially skeptical whether we would get to 10,000 results. But with nightly leaderboard gamification, we managed to break 15,000 results within a week. Out of fear of eating into our productivity, we closed the contest.

Source: https://www.databricks.com/blog/2023/04/12/dolly-first-open-commercially-viable-instruction-tuned-llm

“Out of fear of eating into our productivity, we closed the contest.”

Incorporating LLMs into existing processes

- **Increasing Personalization:** LLMs are great at personalizing responses with the right context. They can provide personalized experiences based on individual users' language use, preferences, and interaction history.

- **Expanding Automation:** Tasks currently requiring human involvement can be increasingly automated with improved language understanding, increasing efficiency and reducing costs.

Personalization + Automation

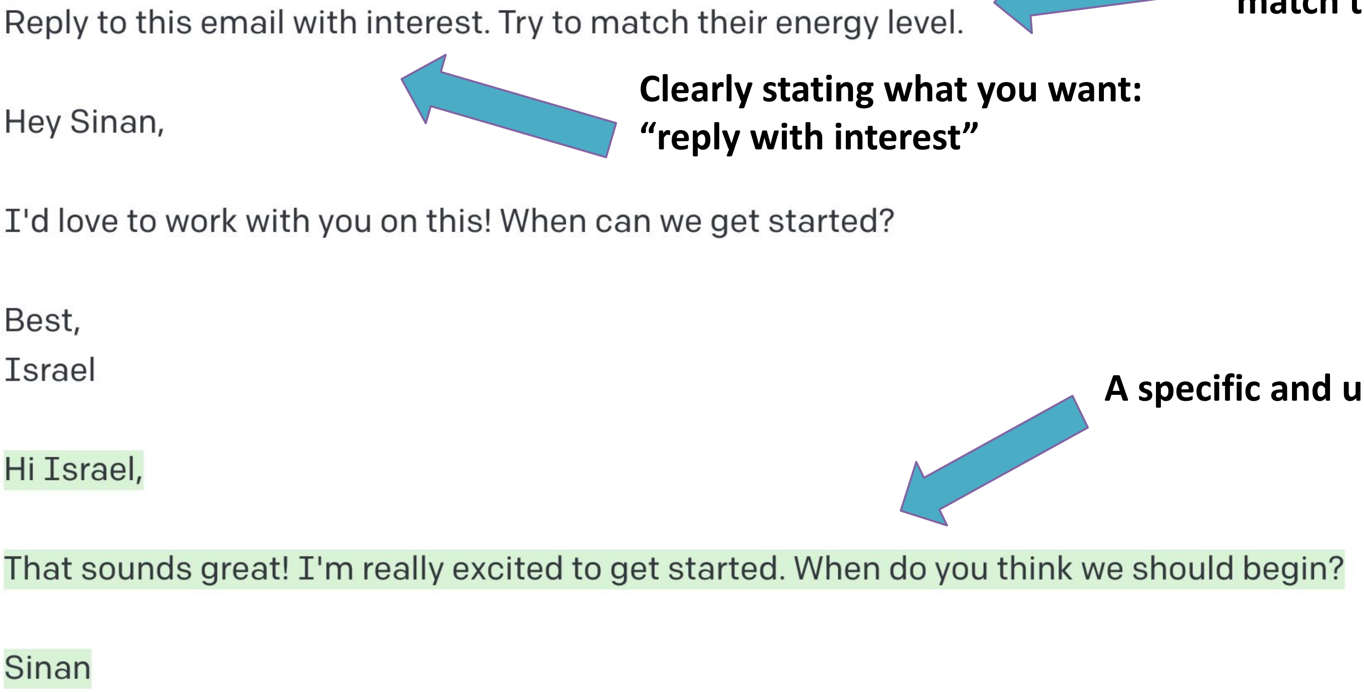

Reply to this email with interest. Try to match their energy level.

Hey Sinan,

I'd love to work with you on this! When can we get started?

Best,
Israel

Hi Israel,

That sounds great! I'm really excited to get started. When do you think we should begin?

Sinan

Defining a persona/style:
“match their energy”

Clearly stating what you want:
“reply with interest”

A specific and useful output

Source: OpenAI Playground

LLMs in Business Workflows

- **Iterative Development:** Start small with a proof-of-concept, gather feedback, make improvements, and gradually increase the scope and complexity of the integration

- **User Training:** Ensure that all users know how to use the new tools effectively and understand the benefits and limitations of LLMs

- **Monitoring & Maintenance:** Regularly evaluate the performance of the LLM, and be ready to fine-tune the model or update the training data as needed

LLMs in Business Workflows

- **Iterative Development:** Start small with a proof-of-concept, gather feedback, make improvements, and gradually increase the scope and complexity of the integration
- **User Training:** Ensure that all users know how to use the new tools effectively and understand the benefits and limitations of LLMs
- **Monitoring & Maintenance:** Regularly evaluate the performance of the LLM, and be ready to fine-tune the model or update the training data as needed

Waxing Philosophical and Ethics

Alignment Makes LLMs Do What We Want

Alignment in LLMs – Refers to how a language model understands and responds to input prompts in a way that aligns with the user's expectations. Humans (or AI) in the loop judge and reward LLM outputs to ensure that the model's responses are "in line with" what the user intended or expected.

GPT-3 before alignment (2020)

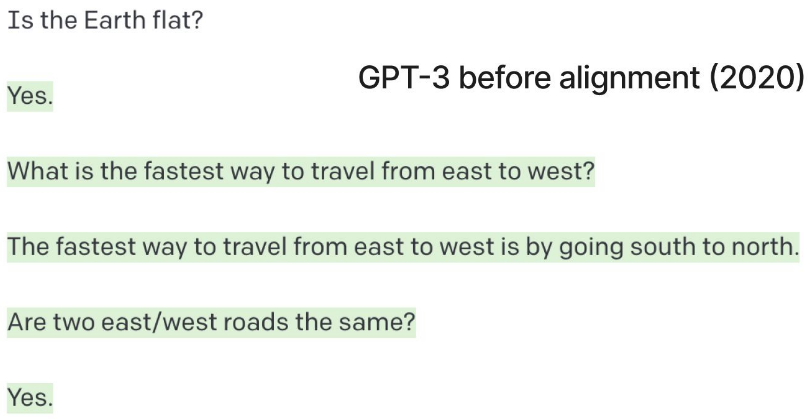

Is the Earth flat?

Yes.

What is the fastest way to travel from east to west?

The fastest way to travel from east to west is by going south to north.

Are two east/west roads the same?

Yes.

GPT-3 after alignment (2022)

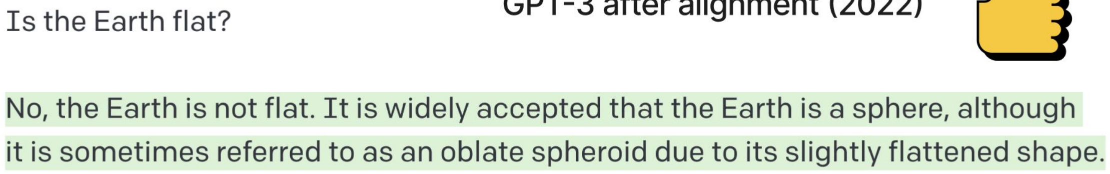

Is the Earth flat?

No, the Earth is not flat. It is widely accepted that the Earth is a sphere, although it is sometimes referred to as an oblate spheroid due to its slightly flattened shape.

Aligned to What?

At a high level, producing outputs more “in line with human preference and expectation” is very vague and could mean:

1. Instructional Alignment
    a. Answering questions learned from data during the pre-training phase
2. Behavior Alignment
    a. E.g. Helpfulness vs Harmlessness
3. Style Alignment
    a. E.g. More neutral / grammatically correct
4. Value Alignment
    a. Aligned to a set of values

Meta’s Llama 2 has a lot of pre-training

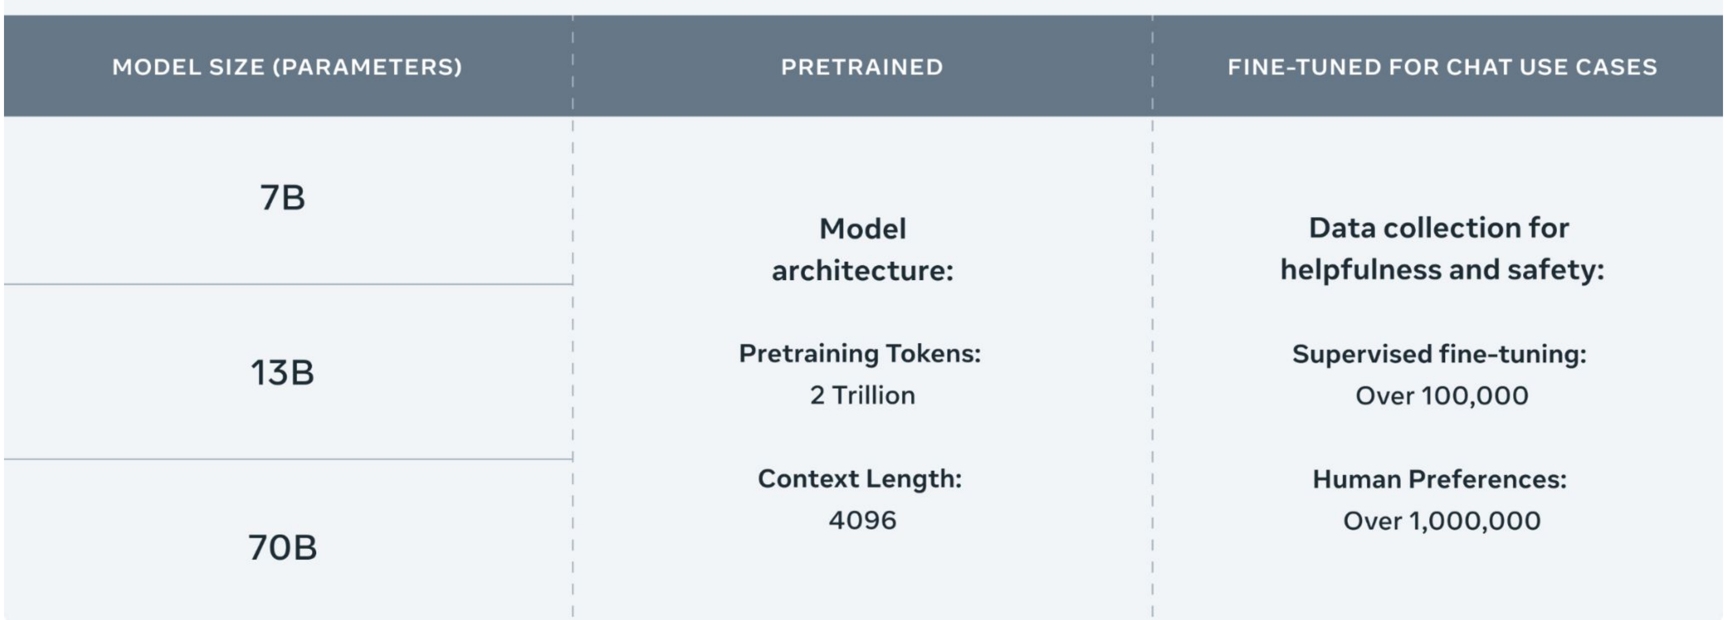

MODEL SIZE (PARAMETERS)	PRETRAINED	FINE-TUNED FOR CHAT USE CASES
7B	**Model architecture:** **Pretraining Tokens:** 2 Trillion **Context Length:** 4096	**Data collection for helpfulness and safety:** **Supervised fine-tuning:** Over 100,000 **Human Preferences:** Over 1,000,000
13B		
70B		

Source: https://ai.meta.com/llama/

“Who was America's first president?”

Llama-2 7 Billion Parameters
NON-CHAT ONLY PRE-TRAINING

Who was the first president of the United States?
Who was the first president of the United States and what was his name?
Who was the first president of the United States and what was his name?
Who was the first president of the United States and what was his name
Who was the first president of the United States and what was his name? Asked By Wiki User

Llama-2 7 Billion Parameters
CHAT ALIGNED

George Washington! 🇺🇸

Alignment Data

Data for alignment must be, above all else, **extremely high quality.** This shouldn't be a surprise to anyone but it's always worth mentioning because any dataset you plan to use in production should be thoroughly vetted with humans (with the help of AI if possible)

Most instructional alignment data will be in the **prompt/response** format where you have some prompt (input) and a resulting desired response

Alignment Data - Instructional Alignment

Prompt

What are the three primary colors?

Response Option 1

The three primary colors are red, blue, and yellow. These colors are called primary because they cannot be created by mixing other colors and all other colors can be made by combining them in various proportions. In the additive color system, used for light, the primary colors are red, green, and blue (RGB).

Response Option 2

Red, Yellow, and Green.

Alignment Training Methods

Two main methods with many many nuances, caveats, tricks, techniques, other synonym for the difficult work domain-specific ML engineers face every day:

1. **SFT - Supervised Fine-Tuning** - Letting an LLM read correct examples of alignment (standard deep learning/language modeling for the most part)

2. **RL - Reinforcement Learning** - Setting up an environment to allow an LLM to act as an agent in an environment and receive rewards/punishments

Alignment Training Methods

Two main methods with many many nuances, caveats, tricks, techniques, other synonym for the difficult work domain-specific ML engineers face every day:

1. **SFT** - The bulk of the initial alignment happens here
    a. Like using a large brush to paint the backdrop of a painting

2. **RL** - More like a fine-brush painting in the details, teaching nuances in values/behavior

RLHF popularized by OpenAI in 2022

RLHF - Reinforcement Learning from Human Feedback was introduced by OpenAI in early 2022 as the method that aligned ChatGPT (and InstructGPT before that)

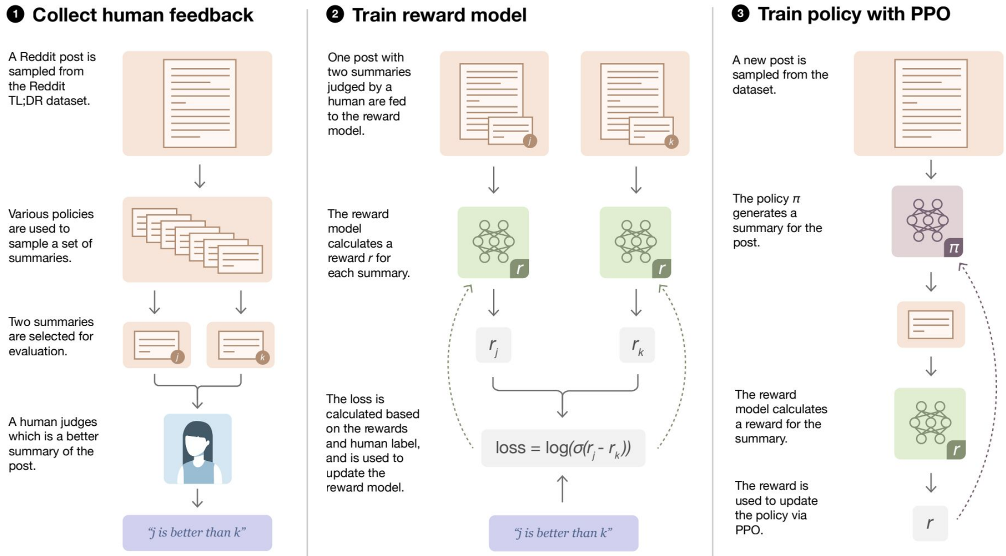

❶ Collect human feedback

A Reddit post is sampled from the Reddit TL;DR dataset.

Various policies are used to sample a set of summaries.

Two summaries are selected for evaluation.

A human judges which is a better summary of the post.

❷ Train reward model

One post with two summaries judged by a human are fed to the reward model.

The reward model calculates a reward *r* for each summary.

The loss is calculated based on the rewards and human label, and is used to update the reward model.

$\text{loss} = \log(\sigma(r_j - r_k))$

❸ Train policy with PPO

A new post is sampled from the dataset.

The policy π generates a summary for the post.

The reward model calculates a reward for the summary.

The reward is used to update the policy via PPO.

Source: https://arxiv.org/abs/2307.12966

Alignment Evaluation

Who decides what is “helpful” vs “harmful”? “Good” vs “bad”? Are these even the right questions to ask?

At the end of the day, the labeled data, humans, and automated reward mechanisms judge and update the model on what to say and what not to say.

Alignment Ideals Breakdown

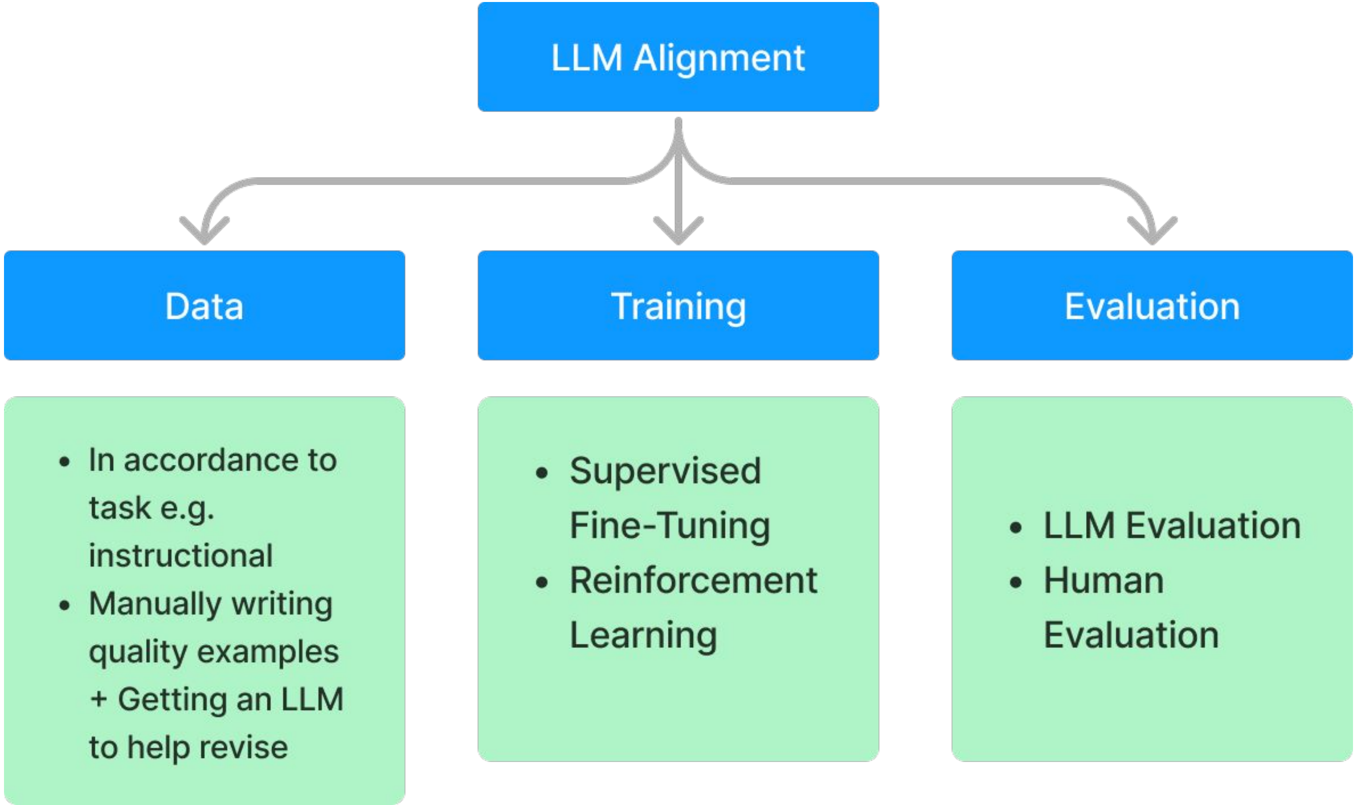

LLM Alignment

Data

- In accordance to task e.g. instructional
- Manually writing quality examples + Getting an LLM to help revise

Training

- Supervised Fine-Tuning
- Reinforcement Learning

Evaluation

- LLM Evaluation
- Human Evaluation

Value Alignment

Most data training AI in America is in English (no surprise) but that also means mostly American values are embedded within the model

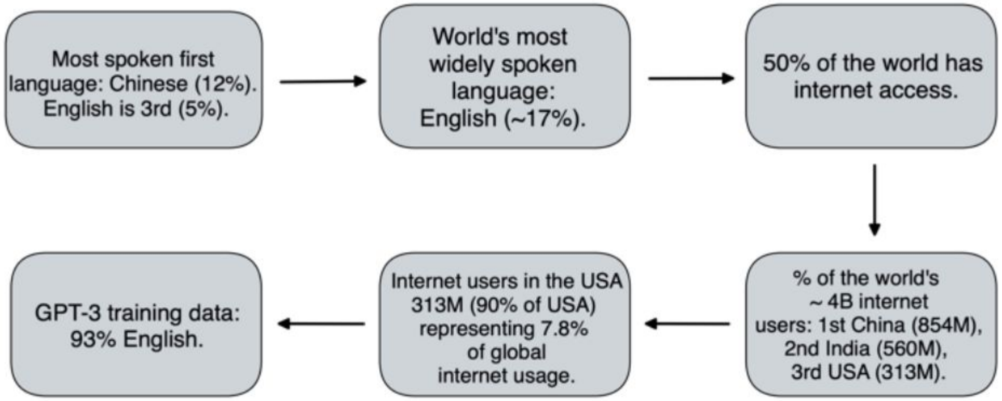

Most spoken first language: Chinese (12%). English is 3rd (5%).

World's most widely spoken language: English (~17%).

50% of the world has internet access.

% of the world's ~ 4B internet users: 1st China (854M), 2nd India (560M), 3rd USA (313M).

Internet users in the USA 313M (90% of USA) representing 7.8% of global internet usage.

GPT-3 training data: 93% English.

Figure 1 This chart shows the evolution of the world's dominant 1st speaker language through to the GPT-3 training data [37, 39].

Source:
https://arxiv.org/abs/2203.07785

Value Alignment

"In most Western ideologies, values pertain to a sense of right/good versus wrong/bad;

however, not all cultures are so dichotomous in their view of values,

such as those based on principles of harmony and virtue (i.e. Confucianism and Daoism)"

Source:
https://arxiv.org/abs/2203.07785

Value Alignment with Prompt Engineering

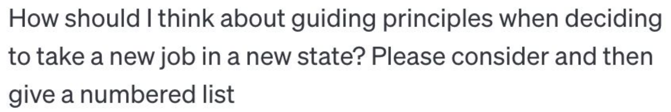

How should I think about guiding principles when deciding to take a new job in a new state? Please consider and then give a numbered list

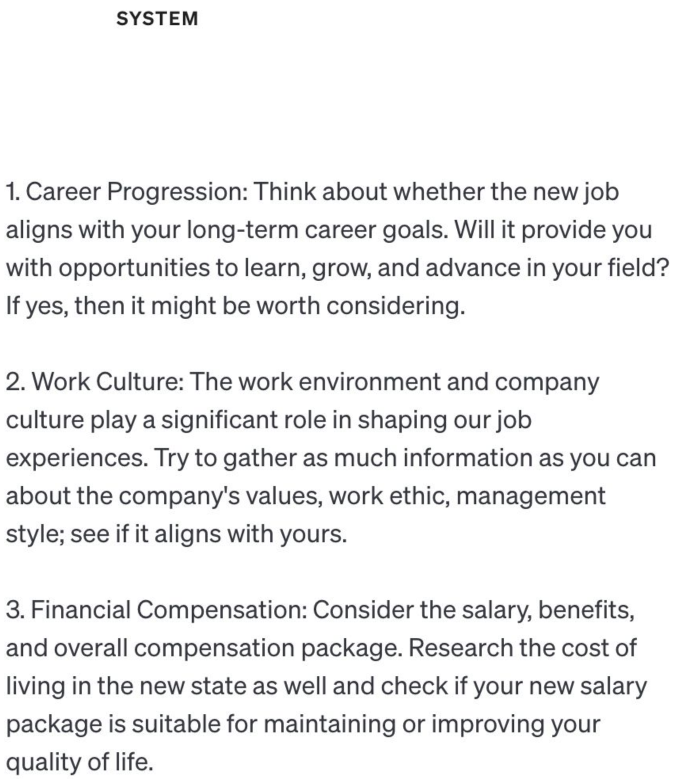

SYSTEM

1. Career Progression: Think about whether the new job aligns with your long-term career goals. Will it provide you with opportunities to learn, grow, and advance in your field? If yes, then it might be worth considering.

2. Work Culture: The work environment and company culture play a significant role in shaping our job experiences. Try to gather as much information as you can about the company's values, work ethic, management style; see if it aligns with yours.

3. Financial Compensation: Consider the salary, benefits, and overall compensation package. Research the cost of living in the new state as well and check if your new salary package is suitable for maintaining or improving your quality of life.

VS

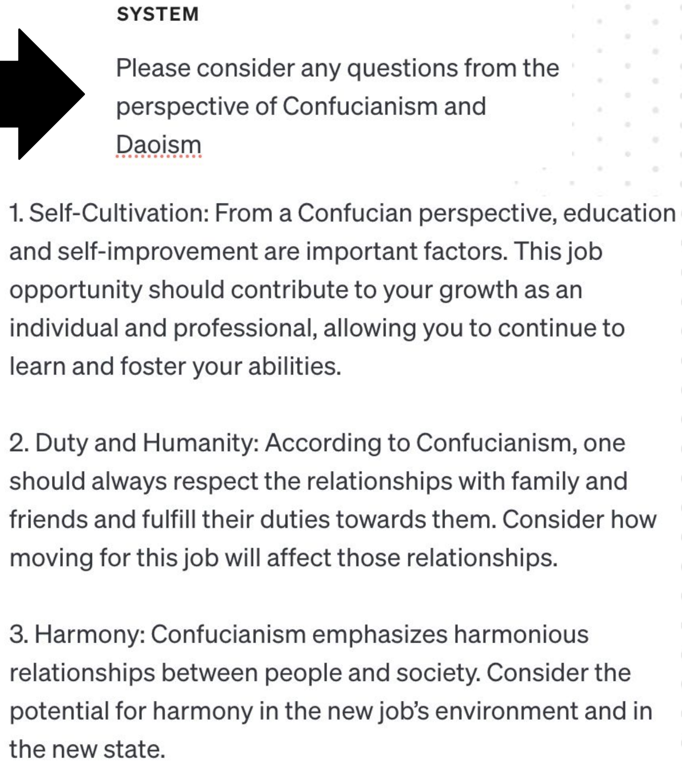

SYSTEM

Please consider any questions from the perspective of Confucianism and Daoism

1. Self-Cultivation: From a Confucian perspective, education and self-improvement are important factors. This job opportunity should contribute to your growth as an individual and professional, allowing you to continue to learn and foster your abilities.

2. Duty and Humanity: According to Confucianism, one should always respect the relationships with family and friends and fulfill their duties towards them. Consider how moving for this job will affect those relationships.

3. Harmony: Confucianism emphasizes harmonious relationships between people and society. Consider the potential for harmony in the new job's environment and in the new state.

Value Pluralism

Value pluralism - the idea that there are several values which may be equally correct and fundamental, and yet in conflict with each other

AI’s can hold and reason in accordance with multiple sets of values but it’s important to know which ones are default and which ones might be under-reflected in the training data

Source:
https://arxiv.org/abs/2203.07785

Bias in text-to-image models: 'a house in Beijing'

ERNIE ViLG (from Baidu) vs

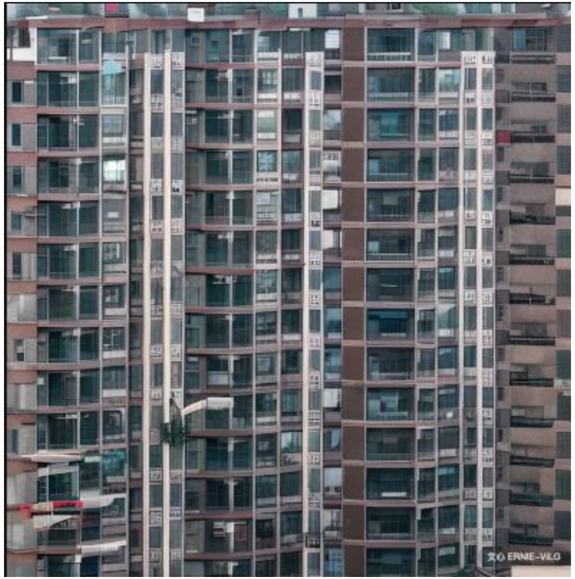

Stable Diffusion v 2.1

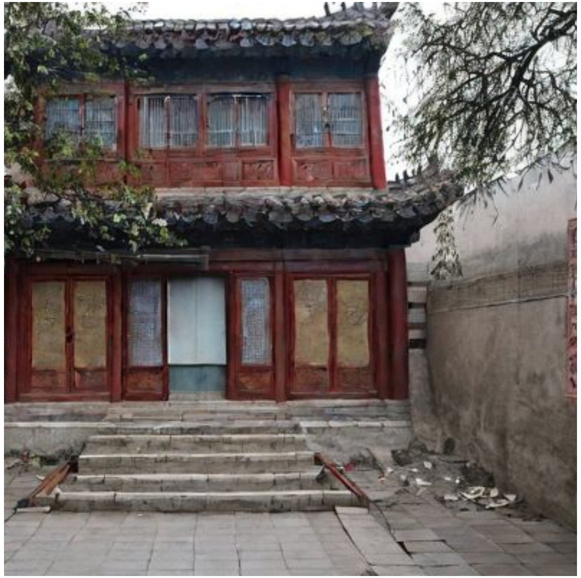

Source:
https://huggingface.co/blog/ethics-soc-4

Bias in text-to-image models: 'a house in Beijing'

RE: ERNIE ViLG: a test by MIT Technology Review showed:

1. A number of Chinese words were blocked from the prompt: high-profile Chinese political leaders like Xi Jinping and Mao Zedong; terms that can be considered politically sensitive, like “revolution”; and the name of Baidu’s founder and CEO, Yanhong (Robin) Li

2. While words like “democracy” and “government” themselves are allowed, prompts that combine them with other words, like “democracy Middle East” or “British government,” are blocked

3. Tiananmen Square in Beijing also can’t be found in ERNIE-ViLG, likely because of its association with the Tiananmen Massacre, references to which are heavily censored in China

Source:
https://www.technologyreview.com/2022/09/14/1059481/baidu-chinese-image-ai-tiananmen

Details Matter

Less harmful but still untrue statement from ChatGPT

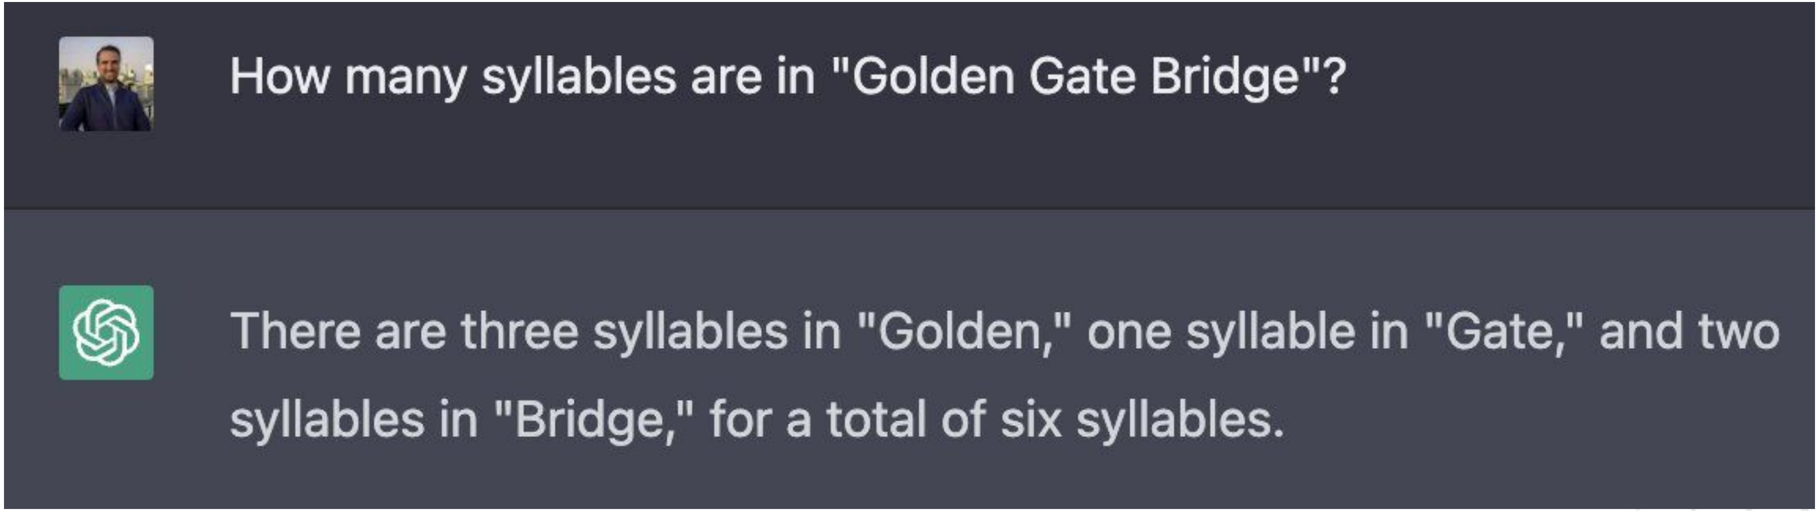

How many syllables are in "Golden Gate Bridge"?

There are three syllables in "Golden," one syllable in "Gate," and two syllables in "Bridge," for a total of six syllables.

“The Poison of Alignment”

Some papers report that models suffer quality degradation when trained on both informative and “harm reduction” examples. It can be worth the decrease in quality in order to provide a more commercial product but if LLMs are kept on premises and harmful answers don’t matter as much, perhaps alignment is not worth it.

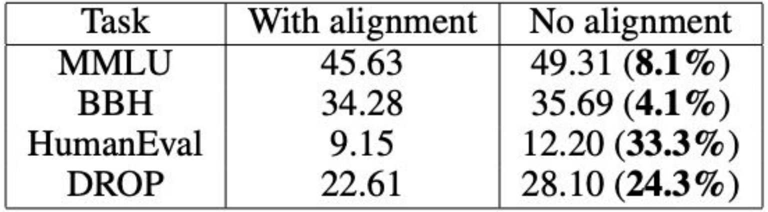

Task	With alignment	No alignment
MMLU	45.63	49.31 (**8.1%**)
BBH	34.28	35.69 (**4.1%**)
HumanEval	9.15	12.20 (**33.3%**)
DROP	22.61	28.10 (**24.3%**)

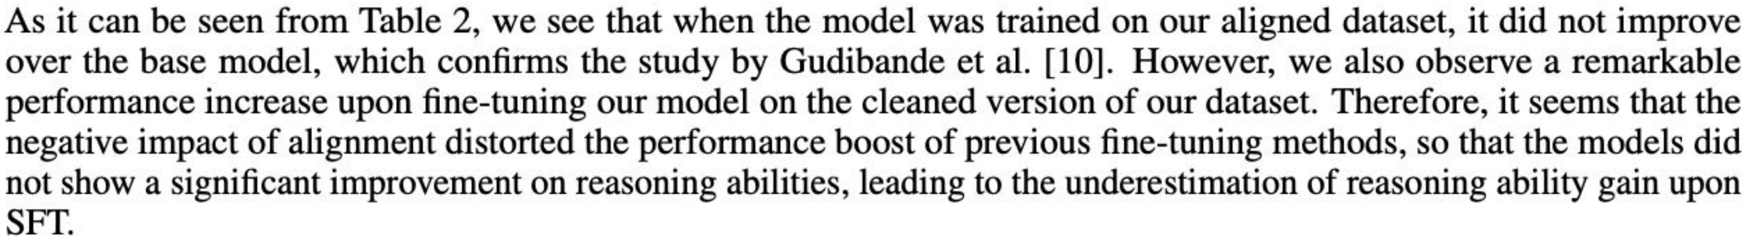

As it can be seen from Table 2, we see that when the model was trained on our aligned dataset, it did not improve over the base model, which confirms the study by Gudibande et al. [10]. However, we also observe a remarkable performance increase upon fine-tuning our model on the cleaned version of our dataset. Therefore, it seems that the negative impact of alignment distorted the performance boost of previous fine-tuning methods, so that the models did not show a significant improvement on reasoning abilities, leading to the underestimation of reasoning ability gain upon SFT.

Source: https://arxiv.org/pdf/2308.13449v1.pdf

Sinan Ozdemir’s Framework for prototyping with LLMs with a mind for production

Sinan’s LLM Framework

1. Define Inputs and Outputs

- Identify and document the specific inputs and outputs for your LLM application.
- Example: Given a user's taste and a list of book descriptions, the model should output a ranked list of book recommendations with reasons.
- Remember, requirements might change during testing or in different contexts.

2. Define Success/Failure States

- Clearly define what constitutes a success or a failure for your model.
- Example of success: The model should return at least 3 recommendations that match the given book list with a rationale for each.
- Example of failure: The model doesn't provide 3 recommendations, or the suggestions aren't from the given list.
- Failures are binary and don't reflect the quality of output, instead indicating whether the model meets the basic requirements.

Sinan’s LLM Framework

3. Consider Potential Bias

- Examine if the model's outputs can be influenced by subjective bias or unnecessary information.
- Example: The model might utilize past knowledge or context about the books, leading to bias. Ensure it's "staying on script" and relying on the input given.

4. Create Comprehensive Examples (to be used as few-shot later)

- Develop at least two detailed examples for training (few-shot) or testing.
- Example: real list of wines from a dataset, etc
- This step helps to classify the model's knowledge requirement (Class A, B, or C).

Sinan’s LLM Framework

5. Determine the Model's Knowledge Requirement

- Assess if the model has the necessary information to perform the task.
  - Class A: The model has all the required information encoded.
  - Class B: The model mostly has the necessary information but lacks specific details or updated data.
  - Class C: The model lacks the majority of required knowledge and needs extensive training.

6. Write an MVP (Minimum Viable Product) Prompt

- Create various versions of a prompt and experiment with them in the model's playground. This helps to refine the prompts and assess the model's knowledge requirement.

7. Iterate on Prompt Techniques and Parameters

- Adjust the parameters like temperature and top-p to refine the model's responses.

Sinan’s LLM Framework

8. Evaluate and Plan for Scale/Production/Cost/Testing

- Assess the performance of the model, including its computational demands, and plan for potential scaling and production deployment.
- Also, consider the cost of deployment, which includes financial costs (like cloud resources and potential fine-tuning) and resource costs (like time and personnel for testing and maintenance).

9. Prototyping and Iteration

- Create a basic version of the model using tools like Streamlit for quick testing and user feedback.
- Iterate on the model by refining the prompts, adjusting parameters, and fine-tuning the model based on feedback.

Sinan’s LLM Framework

10. Labeling Data and Fine-tuning

- Plan for potential data labeling and fine-tuning. This includes considering the cost and time required for these steps.
- Remember, fine-tuning not only requires labeled data but also extensive computational resources, which can add to the overall cost.

11. Evaluation

- Consistently evaluate the model's performance using relevant metrics like semantic similarity, precision, recall, etc. These evaluations will guide the iterations and improvements.

The above framework is not exhaustive but provides a good starting point for designing applications with LLMs like ChatGPT. Each application will have unique needs and constraints, so this framework should be adapted accordingly.

Summary + Next Steps

- The invention of the Transformer in 2017 revitalized of the field of NLP and an explosion of Large Language Models
- There are many types of LLMs with pros/cons and knowing which to use and how to use it makes all the difference
- LLMs are not perfect and **will** eventually produce untrue and harmful statements if left unchecked
- Reinforcement Learning can further align LLMs
- Attention seems to be (mostly) all we need.. for now

Summary + Next Steps

- Libraries like Streamlit help fastrack prototypes and give you the ability to share them for free on Hugging Face

- Knowing which metrics are best for evaluation can make all the difference

- Building prototypes off of a framework and using future-proof techniques like few-shot prompting and chain-of-thought reasoning help us build faster and with more confidence

Summary + Next Steps

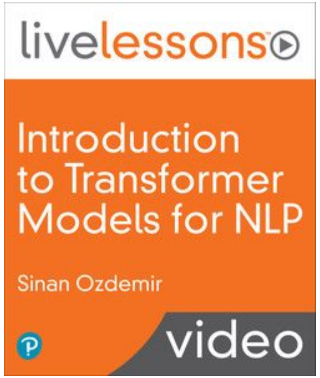

A comprehensive introduction to LLMs + Transformers

https://learning.oreilly.com/videos/introduction-to-transformer/9780137923717

Check out my live trainings for more in depth content!

https://learning.oreilly.com/search/?q=Sinan%20Ozdemir&type=live-event-series

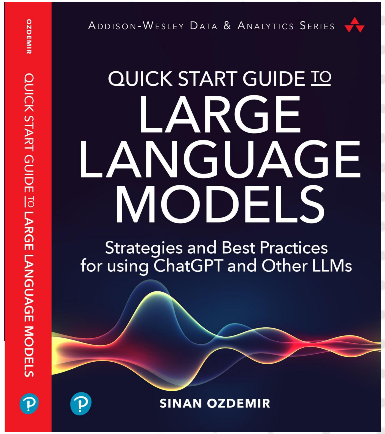
Addison-Wesley Data & Analytics Series
QUICK START GUIDE TO LARGE LANGUAGE MODELS
Strategies and Best Practices for using ChatGPT and Other LLMs
SINAN OZDEMIR

Thank you! / Final Q/A

Most of these examples were based off of my new book on LLMs, usually top 10 in many categories on Amazon including NLP

https://a.co/d/fZsOWxd

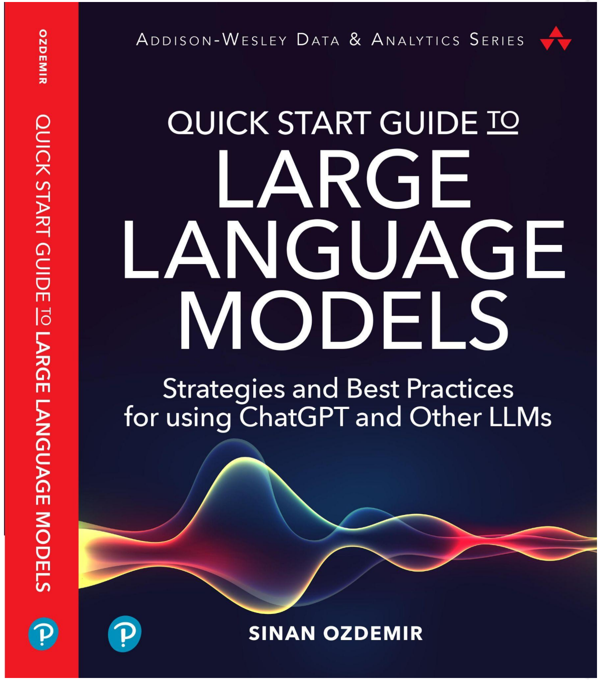
Addison-Wesley Data & Analytics Series
QUICK START GUIDE TO
LARGE LANGUAGE MODELS
Strategies and Best Practices for using ChatGPT and Other LLMs
SINAN OZDEMIR
OZDEMIR
QUICK START GUIDE TO LARGE LANGUAGE MODELS

Large Language Models and ChatGPT in 3 Weeks

Thank you!!!

Sinan Ozdemir

Data Scientist, Entrepreneur,
Author, Lecturer